

# Parallel Design Patterns-L01

An introduction to the course and  
parallel design patterns

---

Course Organiser: Dr Nick Brown  
nick.brown@ed.ac.uk  
Bayes room 2.13

- Lectures:
  - Tuesdays 15:10 – 16:00 in LG.06: 40 George Square
  - Wednesdays 12:10 – 13:00 in LG.06: 40 George Square
- Practicals
  - Mondays 11:10 – 12:00 (*starting next week, week two*)
    - *In Appleton tower 6.06*
- Assessment: 100% by coursework, split in two
  - **Part One** (Report; 30%)
    - Handed out towards end of week 3
    - Deadline: 14<sup>th</sup> February, 2025 at 4pm
  - **Part Two** (Code & Report; 70%)
    - Handed out around week 7
    - Deadline: 28<sup>th</sup> March, 2025 at 4pm



*Caoimhin Laoide-Kemp will also be teaching lectures and practicals*



- This is a more abstract course than some others, but will lead to a practical piece of coursework and will revisit ideas you have already met in practice in Semester 1.
- So far most courses have taken a bottom-up approach
  - This course will now look at things from the top, down
- Two important ideas
  - Reusable patterns
  - Parallel frameworks
- Typically look at 1 or 2 patterns per lecture
  - Abstractly describe and relate to languages, hardware and applications
  - Practicals look at implementing patterns to solve a problem
  - Coursework: Identifying relevant patterns, and writing code to implement and apply a parallel framework


## *Patterns for Parallel Programming*

**Mattson, Sanders, Massingill**

Addison Wesley (2005)

*ISBN-10: 0321228111*

*ISBN-13: 978-0321228116*



- The closest text to this course
- Covers the same patterns, generally uses the same terms
- Examinable material is covered in lectures, but if you're a book person, this is the one I'd recommend.

# Planned Course Timetable 2025

Semester Week No	Week Beginning	Lectures	Topics
W1	13 <sup>th</sup> January, 2025	L1, L2	Introduction, Parallel algorithm analysis, Finding Concurrency
W2	20 <sup>th</sup> January, 2025	L3, P1, L4	Geometric Decomposition, Recursive Data, Task Parallelism, Divide & Conquer
W3	27 <sup>th</sup> January, 2025	L5, P2, L6	Pipelines, Event-Based Coordination, Actors <b>Coursework 1 Handed Out</b> (29 <sup>th</sup> January)
W4	3 <sup>rd</sup> February, 2025	L7, P3, L8	Parallel framework design
W5	10 <sup>th</sup> February, 2025	L9, P4, L10	Implementation Strategies, SPMD, Master/Worker, Loop Parallelism, Shared Data <b>Coursework 1 Due: Friday 14<sup>th</sup> February 4pm</b>
ILW	17 <sup>th</sup> February, 2025	No lectures	<b>Flexible Learning Week</b>
W6	24 <sup>th</sup> February, 2025	L11, P5, L12	<b>Coursework 2 Handed Out</b> (26 <sup>th</sup> February) Distributed Array, Fork/Join, Shared Queue
W7	3 <sup>rd</sup> March, 2025	L13, P6, L14	Active messaging, In-situ data analytics, Vectorisation, Memory bound codes, spatial compute
W8	10 <sup>th</sup> March, 2025	L15, P7, L16	High performance data analytics, Domain Specific Languages (DSLs), fosters Methodology
W9	17 <sup>th</sup> March, 2025	No lectures	<i>Practicals as an open surgery for coursework</i>
W10	24 <sup>th</sup> March, 2025	No lectures	<i>Practicals as an open surgery for coursework</i>
W11	31 <sup>st</sup> March, 2025	No lectures	<b>Coursework 2 Due: Friday 28<sup>th</sup> March, 4pm</b>

Term	Description
Task	Sequence of instructions that operate together as a group which corresponds to some logical part of the code.
Unit of Execution (UE)	To be executed a task needs to be mapped to a unit of execution – such as a process or a thread. This is a generic term for a collection of possibly concurrent executing entities
Processing Element (PE)	Some hardware element to execute the UEs. A single SMP machine might be one PE, whereas in a distributed machine (such as ARCHER) a PE would be a node.

- Motivation: The same concepts and problem types appear in many different places
- We don't want to waste time re-inventing the wheel
- We'd like a common language to talk about “ways of doing parallelism” between different, non HPC expert, stake holders
- Languages, machines and applications change frequently but ideas and concepts recur
- We often start with some problem/code in an area we know little about. Can help us know where to start.

# What is a Design Pattern?

- The idea of a design pattern was first formally described by the architect Christopher Alexander in the field of architecture in his 1977 book\*
- “Each pattern describes a problem that occurs over and over again in our environment, and then describes the core of the solution to that problem, in such a way that you can use this solution a million times over, without ever doing it the same way twice” – Christopher Alexander

\* Alexander, Christopher (1977). A Pattern Language: Towns, Buildings, Construction. Oxford University Press. ISBN 0195019199



- Sharing  $n$  things of type  $t$  amongst  $m$  people
  - Doesn't matter what  $n$ ,  $t$ , and  $m$  are
- Sorting algorithms
  - As long as you have an ordering amongst any two items, you can use the same algorithm to sort strings, numbers, whatever.

# What is a Design Pattern?

---

- A description of a problem and a strategy for its solution expressed in an abstract way independent of language, hardware, and application
- “A design pattern describes a good solution to a recurring problem in a particular context” – *Mattson et al*
- “a design pattern is a general reusable solution to a commonly occurring problem within a given context” – *Wikipedia*

- First example of Design Patterns used in software engineering: Beck & Cunningham (1987)
- Design Patterns in the field of software engineering popularised by the “gang of four”:
  - Erich Gamma, Richard Helm, Ralph Johnson and John Vlissides
- *Design Patterns: Elements of Reusable Object-Oriented Software* (Addison-Wesley, 1995. ISBN 0-201-63361-2)

- The GoF didn't come up with the ideas described in these patterns, but they were the first to categorize them and describe them as patterns
- Creational Patterns, e.g.,
  - Factory method, Singleton
- Structural Patterns, e.g.,
  - Decorator, Façade, Proxy
- Behavioural Patterns, e.g.,
  - Iterator, Visitor, Observer (often used as part of Model-View-Controller)

**We are not going to study the Gang of Four's Design Patterns!**



- These are **design** patterns because they are used during the design of a piece of software or a system
- They should help you to think about a solution to a problem before any implementation in code
- They are **not a process**
- There is rarely *one right answer* and a good design often boils down to a number of *tradeoffs*

# Patterns in a Design Process

An example from *Patterns for Parallel Programming*<sup>1</sup>

## Finding Concurrency

- Task Decomposition, Data Decomposition, Group Tasks, Order Tasks, ...

## Algorithm Structure

- Tasks Parallelism, Divide and Conquer, Geometric Decomposition, Recursive Data, ...

## Supporting Structures

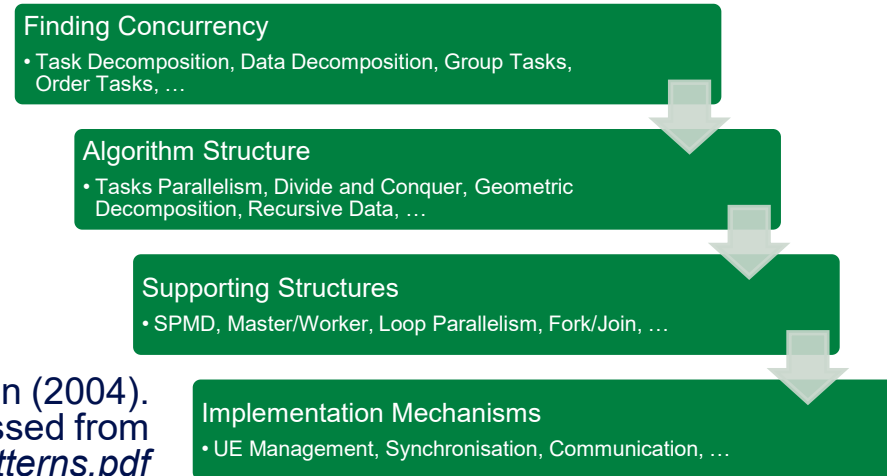
- SPMD, Master/Worker, Loop Parallelism, Fork/Join, ...

## Implementation Mechanisms

- UE Management, Synchronisation, Communication, ...

<sup>1</sup> Patterns for Parallel Programming; Mattson, Sanders, Massingill; Addison Wesley (2005)

- Patterns can be grouped into “Strategies” or “Design Spaces”
- The grouping is sometimes referred to as a Pattern Language
  - “Pattern Language - a collection of design patterns, guiding users through the decision process in building a system” – Kim (2004) \*
- Parallel Algorithm Strategy
  - *aka* “Algorithm Structure Design Space”
- Implementation Strategy
  - *aka* “Supporting Structure Design Space”
  - distinct from “Implementation Mechanisms Design Space”



\* Eun-Gyu Kim, *Parallel Patterns*, online presentation (2004).  
Accessed from  
<http://web.engr.illinois.edu/~snir/patterns/patterns.pdf>

- On algorithm structure and supporting structures
- Discussion about *finding concurrency* next lecture

## Finding Concurrency

- Task Decomposition, Data Decomposition, Group Tasks, Order Tasks, ...

## Algorithm Structure

- Tasks Parallelism, Divide and Conquer, Geometric Decomposition, Recursive Data, ...

## Supporting Structures

- SPMD, Master/Worker, Loop Parallelism, Fork/Join, ...

## Implementation Mechanisms

- UE Management, Synchronisation, Communication, ...

- Implementation mechanisms dealt with elsewhere
  - Will use implementation technologies (MPI and OpenMP) in the practicals
  - Details of how hardware, operating system and middleware can implement the parallel algorithm at run-time
  - Covered in other modules
    - Thread and Process management
    - Low-level synchronisation mechanisms
    - Communications constructs and protocols



# Patterns in a Design Process

An example from *Patterns for Parallel Programming*<sup>1</sup>

## Finding Concurrency

- Task Decomposition, Data Decomposition, Group Tasks, Order Tasks, ...

## Algorithm Structure

- Tasks Parallelism, Divide and Conquer, Geometric Decomposition, Recursive Data, ...

## Supporting Structures

- SPMD, Master/Worker, Loop Parallelism, Fork/Join, ...

## Implementation Mechanisms

- UE Management, Synchronisation, Communication, ...

<sup>1</sup> Patterns for Parallel Programming; Mattson, Sanders, Massingill; Addison Wesley (2005)

## The *Algorithm Structure* Design Space

- Input information:
    - Grouped set of tasks
    - A knowledge of dependencies amongst tasks and any implied temporal constraints
  - These patterns can be thought of as algorithm templates
- Task Parallelism
  - Divide and conquer
  - Geometric Decomposition (Domain decomposition)
  - Recursive Data
  - Pipelines
  - Event-Based Coordination
  - Actor pattern

# Patterns in a Design Process

An example from *Patterns for Parallel Programming*<sup>1</sup>

## Finding Concurrency

- Task Decomposition, Data Decomposition, Group Tasks, Order Tasks, ...

## Algorithm Structure

- Tasks Parallelism, Divide and Conquer, Geometric Decomposition, Recursive Data, ...

## Supporting Structures

- SPMD, Master/Worker, Loop Parallelism, Fork/Join, ...

## Implementation Mechanisms

- UE Management, Synchronisation, Communication, ...

<sup>1</sup> Patterns for Parallel Programming; Mattson, Sanders, Massingill; Addison Wesley (2005)

## The *Supporting Structures* Design Space

- Usually considered once the Algorithm Structure has been decided
  - Can be divided into *Program Structures* and *Data Structures*
- Master / Worker
  - Loop Parallelism
  - Fork / Join
  - Shared Queue
  - SPMD
  - Shared Data
  - Distributed Array
  - Active messaging
  - Vectorisation



- Pattern Name
  - Should be standard to ease communication
- Problem
- Description of the context
  - A more detailed discussion of where the pattern might be applicable
- The forces
  - Forces that act on the design using a specific pattern
  - Goals and constraints
  - Things that influence whether this is the right pattern to use
- The solution

- We teach you about Parallel Design Patterns because we think they are a useful abstraction, however there are some who criticise design patterns:
- There's nothing new or special about design patterns; they just boil down to reusing an idea and making life easier.
- Writing code to force it to look like a standard pattern can unnecessarily increase complexity
- The “parallel pattern language” is not standardised enough to be useful
  - There are different names for the patterns and strategies

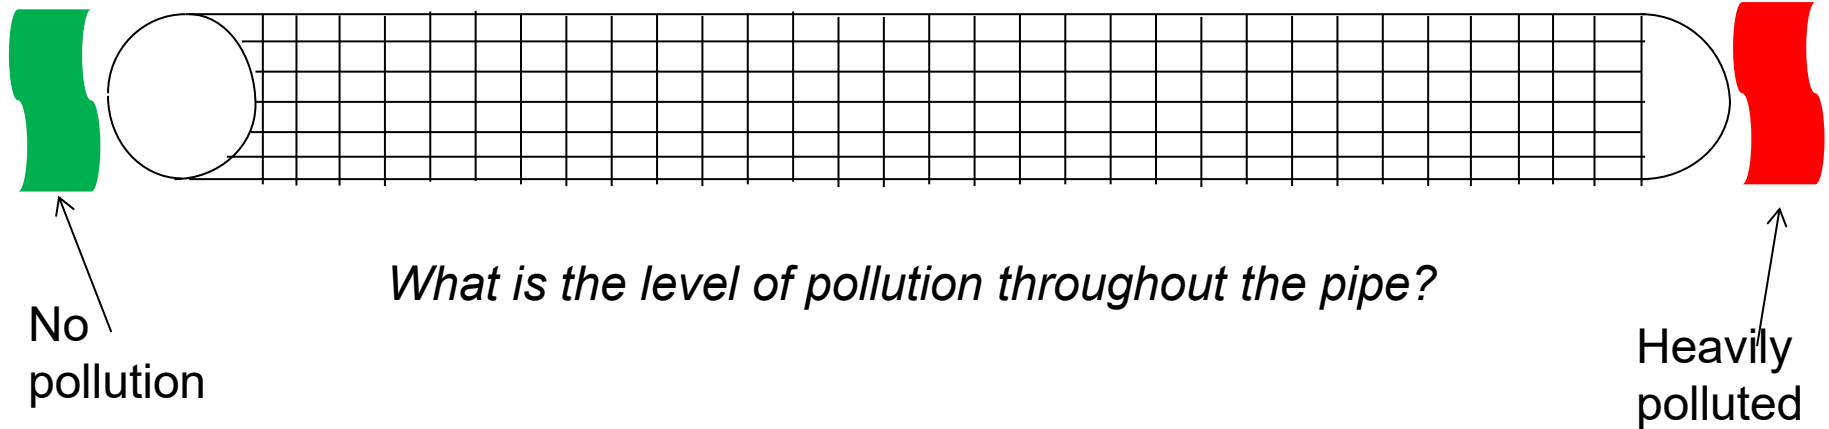
- A design pattern is only necessary because of some missing features of a given programming language
- Parallel Programming “Systems”
  - e.g. DPnDP (Siu, et al)
- Parallel Skeletons
- There are also quite a few alternative approaches
  - e.g. Foster’s Methodology
    - *Partitioning*: Discover as much parallelism as possible
    - *Communication*: Determine communication (local & global)
    - *Agglomeration*: Group individual tasks into larger tasks
    - *Mapping*: Map tasks to UEs

- The other important strand to this course
- A good software engineer will use libraries and frameworks wherever appropriate
- Someone still has to write these libraries and frameworks
- This course aims to introduce some basic ideas about how these might be put together
- Even in a specific code, developing with flexibility and reuse in mind can help significantly when adding additional features.

- We will also look at how HPC might change in the future and consider the impact this will have on how we write parallel codes
  - In terms of changes to hardware forcing us to write applications differently
  - New techniques which are not yet in the mainstream, often build on or complement the well known patterns we will study
  - Some patterns are more mature and/or readily used than others




- **Parallel Design Pattern**
  - A description of a problem (with context) and strategies for its solution, independent of language, architecture and application
  - Not a process
  - Useful, but not a silver bullet
- **Parallel Framework**
  - Ideas from these recurring parallel patterns can be expressed in code to produce parallel frameworks that other programmers can use



- Jacobi iteration solving Laplace's equation for diffusion in 2 dimensions  $\nabla^2 u = 0$

for all grid points

$$u_{\text{new}}(i) = 1/4 * (u[j, i-1] + u[j, i+1] + u[j-1, i] + u[j+1, i])$$



- Works in iterations, solving to a specific residual (accuracy)
  - As we parallelise this, the overall number of iterations and residual should be the same as the serial code which is a nice check

# Overview of serial code

```
double * u_k = malloc(sizeof(double) * (ny+2) * (ny+2)), * u_kp1 = malloc(sizeof(double) * (nx+2) * (ny+2)), *temp;
```

```
initialise(u_k, u_kp1);
```

```
double rnorm=0.0, bnorm=0.0, norm;
```

```
int i, j, k;
```

```
for (i=1;i<=nx;i++) {
```

```
    for (j=1;j<=ny;j++) {
```

```
        bnorm=bnorm+.....
```

```
    }
```

```
}
```

```
bnorm=sqrt(bnorm);
```

```
for (k=0;k<MAX_ITERATIONS;k++) {
```

```
    for (i=1;i<=nx;i++) {
```

```
        for (j=1;j<=ny;j++) {
```

```
            rnorm=rnorm+.....
```

```
        }
```

```
    }
```

```
    norm=sqrt(rnorm)/bnorm;
```

```
    if (norm < CONVERGENCE_ACCURACY) break;
```

```
    for (i=1;i<=nx;i++) {
```

```
        for (j=1;j<=ny;j++) {
```

```
            u_kp1[i]=0.25 * .....
```

```
        }
```

```
    }
```

```
    temp=u_kp1; u_kp1=u_k; u_k=temp;
```

```
    rnorm=0.0;
```

```
}
```


*Sets the pollution values at each end of the pipe and the rest to be zero (initial guess)*

*Compute the initial absolute residual*

*Compute the absolute residual of the current solution, then divide this by bnorm to get the relative residual (how far we have progressed)*

*Termination criteria (level of accuracy met)*

*Jacobi iteration to progress the solution*



- Parallelise it!
  - In 1D using geometric decomposition
  - Start with the simplest approach to halo swapping and then add in extra complexity to optimise this

