

Parallel Design Patterns-L07

Parallel Frameworks I

Course Organiser: Dr Nick Brown
nick.brown@ed.ac.uk
Bayes room 2.13

- Parallel Frameworks
 - Why?
- What makes a good parallel framework?
- Common parallel examples as an illustration of frameworks
 - Common frameworks used in HPC codes

- Lots of people write code
- Quite a lot of people write parallel code
- Few people write good parallel code
- Many parallelisations follow similar patterns
- Is it possible to separate the parallelism from the details of the problem?
 - Not entirely, as you will have already seen and will continue to see in this course
 - Things can be done. There is progress to be made...
 - Not everyone agrees that frameworks are the way to address this problem
- Why reinvent the wheel?

Who will write the parallel frameworks?


- People like you?
- Frameworks have to bridge the gap between applications and the underlying architectures
- Need to abstract away unimportant details of each one from the other *while maintaining as much knowledge as possible about available parallelism*
 - This isn't about writing serial code and getting the compiler/machine to parallelise it; it's about separating out the details of what can be done in parallel from how it is done
 - Continued work required to adapt frameworks for new architectures
 - with the aim of maintaining a stable API
 - Higher level programmers (e.g. scientists) will probably still have to learn to “think parallel”

- An abstraction where software provides some generic functionality that can be used or modified by additional user written code that results in application specific software.
- Reusable semi finished architectures
- E.g. Some combinations of
 - Libraries / APIs
 - Middleware / runtimes
 - Standard ways of expressing parallelism such as skeletons



- The line can be blurred between frameworks and libraries. Frameworks typically manage an aspect of your code (e.g. the parallelism, IO, computation etc..)
- To differentiate consider
 - **Inversion of control:** A framework often controls the program flow in some or all of the application
 - **Extensibility:** A framework should be extensible by the user by them providing some code
 - **Default behaviour:** Sensible defaults are often present if the user does not wish to override this
 - **Immutability:** The framework code itself should remain unchanged but can be extended via the user
 - **Data types:** Frameworks often define their own types which are used in your code
 - **Determine key aspects:** Such as data decomposition

Nine rules for writing parallel codes....




- Written by one of the key people in the UNIX scene
- Whilst this book is about UNIX, the lessons can equally be applied to software more generally
- HPC and UNIX share many similar traits.

<http://catb.org/esr/writings/taoup/html>

- Rule of Modularity
- Rule of Clarity
- Rule of Composition
- Rule of Separation
- Rule of Simplicity
- Rule of Parsimony
- Rule of Transparency
- Rule of Robustness
- Rule of Representation
- Rule of Least Surprise
- Rule of Silence
- Rule of Repair
- Rule of Economy
- Rule of Generation
- Rule of Optimisation
- Rule of Diversity
- Rule of Extensibility

Rule of Modularity

- Write simple parts connected by clean interfaces
- Holds the global complexity down
- Users of your framework can very easily select which bits to use and which bits not to use
 - As opposed to some obscure argument rules in the API which determine this.
 - In HPC want to have a very clear idea of what is running



- Can also be a way in which users can extend your framework (which we will see in an example later)

- Separate policy from mechanism
- Policy
 - Provided by the user (they know what they want best)
- Mechanism
 - Provided by your framework.
 - Often the tricky, crucially important, low level and uninteresting (to a user) aspects such as highly efficient halo swapping.
- Sometimes need careful thought how to inject the policy into the mechanism

- Design for simplicity, add complexity only where you need it
 - Best type of optimisation is not to write it in the first place!
- Pressure to make codes more complex
 - Technical machismo
 - Changing requirements
- Complexity results in bloated, huge buggy, brittle codes
- Actively resist bloat and complexity
 - Small is beautiful
 - Need to accept that there is a high value on simple solutions

- Design for visibility
 - Transparency is when you can look at something and immediately understand what is going on
 - Discoverability is when, with the provided tools, you can investigate something and understand what is going on
- E.g are API calls orthogonal or are there many flags and mode bits which means that a single call might do multiple tasks?
- Is it easy to find the appropriate part of the code called by a specific function?
- What about documentation
 - Tools such as Doxygen

- The child of transparency and simplicity
- Users of a framework might very well use it in ways that the developers have not anticipated
 - But if you say that if something does “X” then it should do “X”
- Unusual parameters
 - Such as extra long or empty strings
 - Negative numbers
- How can you convince potential users that your framework is robust?
 - Good practice such as documentation, unit testing, easy building goes a long way
 - A well thought out design and API

- Your framework should always do the least surprising thing
 - Demands the least amount of learning from the users
 - Most effectively connect to user's pre-existing knowledge
- What is the likely background of your users, are there conventions that you can adopt?
 - For instance from some existing packages they use?
- More general technical tradition
 - E.g. Climate and Forecast (CF) convention for data files
 - Format of configuration files (e.g. FORTRAN NAMELIST)
 - Common command line switches
 - Programming language/implementation technologies

- When your framework has nothing surprising to say, then keep quiet
- User attention is a precious resource
- If you want to have debugging information then have this as a flag
 - Ideally via a preprocessor directive so your code is not checking a conditional at runtime

- Repair what you can, but when you must fail noisily and as soon as possible
- If software can adapt to unexpected conditions then great
 - But some of the worst kind of bugs are when a repair fails and the problem quietly causes some corruption.
- When you must fail, then fail transparently
 - Such as generating a core dump

- Design for the future, because it will be here sooner than you think.
- How extensible is your API?
 - Or would it need to be deprecated?
- Can users easily extend your framework
 - Can new functionality be easily plugged in?
 - Make the joints in your framework flexible
- How extensible are your data formats?
 - Including a version number can go a long way here
 - Self describing data formats can be very useful
- How about machine specific optimisation?


Focusing more on HPC

- These rules are based on Raymond's UNIX experience
- But HPC is a far smaller community with fewer resources
 - People are often working individually on codes
 - Driven by science, frameworks often start life unplanned and grow organically
 - Do you want to be inundated with lots of requests for support after funding has expired?
 - Often a developer/user's task to install required frameworks
- Scalability and performance are crucial
 - Raymond argues that programmer time is more precious than machine time, but in HPC we don't have this luxury
 - In our field performance portability is very important

What makes a good parallel framework?

- There are a number of key objectives
 - **Integration:** With existing codes or codes that do a variety of different functions. Also integration with the community.
 - **Visibility:** Is this framework going to impact upon performance or scalability? Is it bit reproducible & does this matter?
 - **Documentation:** Clear instructions (and ideally examples) about how to use this?
 - **Flexible:** Can I do what I need to using this?
 - **Portable and supported:** Already installed and configured on the systems you plan to use?
 - **Extensible:** Can the framework be extended as your code evolves?

- Portable Extensible Toolkit for Scientific computation
- Used for solving systems of differential equations (such as Laplace's equation from the practical.)
- Geometric decomposition



- Problems plug into the framework, can manage parallelism (if you wish), plenty of different choices such as the solvers and preconditioners to use can be selected at runtime

```
KSP ksp
```

```
DM da
```

```
Vec x
```

```
integer ierr
```

We have already initialised PETSc which will, amongst other things, create its options database

```
call KSPCreate(PETSC_COMM_WORLD,ksp,ierr)
```

```
call KSPSetFromOptions(ksp,ierr)
```

```
call DMDACreate3d(PETSC_COMM_WORLD,DMDA_BOUNDARY_PERIODIC,...,nx,&  
ny,nz,npx,npj,npz,..., da,...)
```

```
call DMSetInitialGuess(da,ComputeInitialGuess,ierr)
```

```
call KSPSetComputeRHS(ksp,ComputeRHS,PETSC_NULL_OBJECT,ierr)
```

```
call KSPSetComputeOperators(ksp,ComputeMatrix,PETSC_NULL_OBJECT,ierr)
```

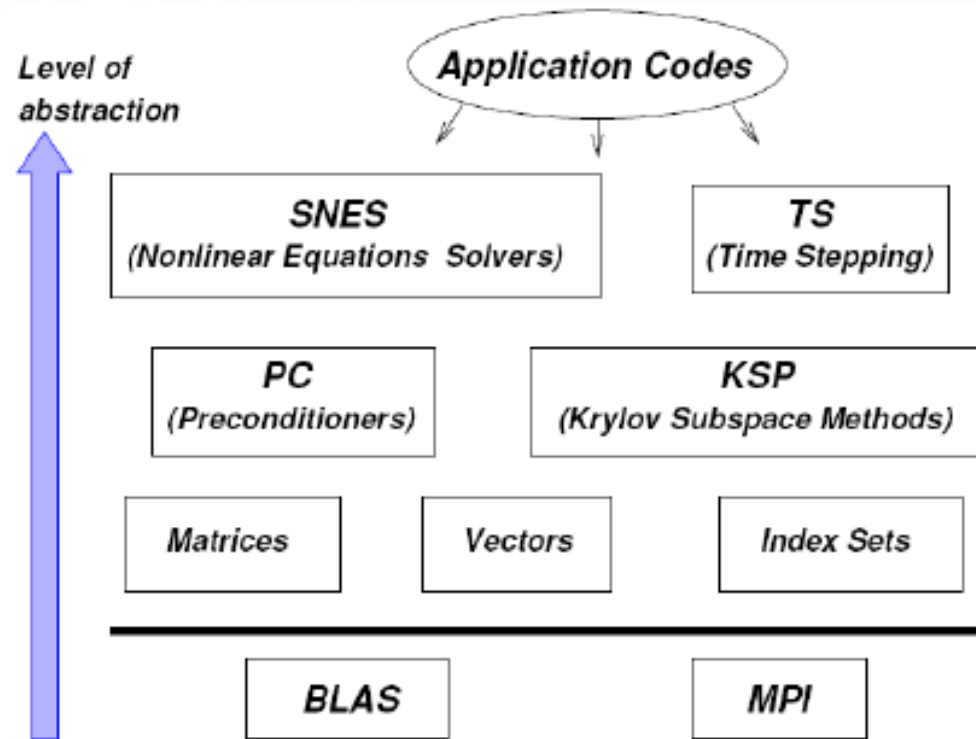
```
call KSPSetDM(ksp,da,ierr)
```

```
call KSPSolve(ksp,PETSC_NULL_OBJECT,PETSC_NULL_OBJECT,ierr)
```

```
call KSPGetSolution(ksp,x,ierr)
```

- Has a beginner, intermediate, advanced and developer API
 - The idea being that it can be understood by a variety of users
 - It is well documented and with plenty of examples
- Supports multiple problem abstractions
 - The scientist's view of $Ax=b$, they provide functions which build up the matrix A and the initial conditions in a vector b . The output from this is the vector x . Using the DM DA calls these can be automatically partitioned and distributed amongst the UEs.
 - Alternatively the programmer can supply routines that define the behaviour which is more advanced but saves explicitly building the matrix
- Numerous third party additions have been written which add additional functionality to the framework

- Operations are defined upon matrices, vectors and index sets, these contain all the parallelism.
 - Sequential matrix and vector results in sequential computation
 - MPI matrix and vector result in MPI parallelisation
 - Threaded matrix and vector result in threaded parallelisation
- In order to support a new architecture (i.e. GPUs or hybrid MPI-OpenMP) then just need to define a new matrix and vector, with the rest of the framework remaining unchanged.



- Facets:
 - **Integration:** Yes, but the temptation can be to rely on the PETSc data structures (such as Mat and Vec)
 - **Visibility:** The solver and preconditioners are well known, published benchmarks for DMDA and is open source.
 - **Flexible:** Lots of options
 - **Portable and supported:** Commonly installed and trivial to build (but does take a long time!) Email support lists
 - **Extensible:** Hierarchical abstraction and organisation
 - **Documentation:** Extensive online documentation and examples. Differentiated for different users and examples
- Disadvantages
 - API can change dramatically between minor releases i.e. 3.5.1 -> 3.5.2 and no backwards compatibility
 - It is possible, but a bit more difficult, to do parallelism without DMDA
 - You might pay a small computation price

- The Fastest Fourier Transformation in the West
- Computing the discrete Fourier transform is a very common operation, hence FFT algorithms are very common



→
e.g. noise removal



“the most important numerical algorithm in our lifetime.” Gilbert Strang, Linear Algebra and Its Applications

- The programmer tells the framework specifics about their problem and this then produces a plan
 - This is (potentially) an expensive operation, so you only want to do it once
 - FFTs are often called in some sort of timestep so you want the actual computation call to be as fast as possible
 - Some FFT kernels also require even data of n^2 , if your data does not look like this then FFT can handle it (or select a different kernel.)
- The programmer then calls FFTW with their plan which will then do the actual FFT.

```
ptrdiff_t N0=n, N1=n
```

```
ptrdiff_t alloc_local, local_n0, local_0_start;
```

```
alloc_local = fftw_mpi_local_size_2d(N0, N1, MPI_COMM_WORLD, &local_n0,  
&local_0_start);
```

```
fftw_complex * data = fftw_alloc_complex(alloc_local);
```

```
fftw_plan plan = fftw_mpi_plan_dft_2d(N0, N1, data, data, MPI_COMM_WORLD,  
    FFTW_FORWARD, FFTW_ESTIMATE);
```

```
... Initialise data to some values ....
```

```
fftw_execute(plan);
```

```
fftw_destroy_plan(plan);
```

- Facets:
 - **Integration:** Yes, easy to limit calls to a specific function. Type definitions are simple and well documented
 - **Visibility:** Benchmarks and source code available for the library
 - **Flexible:** Lots of options and interfaces available
 - **Portable and supported:** Very commonly installed on machines and trivial to compile
 - **Extensible:** Not so much as PETSc, but source code is available
 - **Documentation:** Well documented, lots of examples. API split into beginner, advanced and guru sections.
- Disadvantages
 - It is commonly accepted that the computation aspects are good but the parallel (communication) aspects are not. Therefore often people do their own decompositions and just call the 1D kernels
 - Only slab decomposition in 3D supported, although third party libraries have been developed to address this

- An IO framework, defines and works with files that are self describing and this is very common in HPC scientific applications.



- The basic version is serial, but parallelism is supplied and it is trivial to move to doing parallel IO.
- Uses MPI-IO and HDF5 but the programmer is completely abstracted from the details of these.
- Comes with tools for managing and visualising the data files

Writing a NetCDF file

```
int data_out[NX][NY];  
... populate data_out ....  
  
int ncid, dimids[2], varid;  
  
nc_create(FILENAME, NC_CLOBBER, &ncid);  
nc_def_dim(ncid, "x", NX, &dimids[0]);  
nc_def_dim(ncid, "y", NY, &dimids[1]);  
nc_def_var(ncid, "data", NC_INT, 2, dimids, &varid);  
nc_enddef(ncid);  
  
nc_put_var_int(ncid, varid, &data_out);  
  
nc_close(ncid);
```


Writing a parallel NetCDF file

```
int ncid, dimids[2], varid; size_t start[2], count[2];

start[0] = mpi_rank * NX/mpi_size; start[1] = 0;
count[0] = NX/mpi_size; count[1] = NY;

nc_create_par(FILENAME, NC_NETCDF4|NC_MPIIO, MPI_COMM_WORLD,
              MPI_INFO_NULL, &ncid);

nc_def_dim(ncid, "x", NX, &dimids[0]);
nc_def_dim(ncid, "y", NY, &dimids[1]);
nc_def_var(ncid, "data", NC_INT, 2, dimids, &varid);
nc_enddef(ncid);

nc_put_vara_int(ncid, varid, start, count, &data_out)
nc_close(ncid);
```

Parallel reading a NetCDF file

```
int ncid, varid, size_t start[2], count[2];

nc_open_par(FILENAME, NC_NETCDF4|NC_MPIIO|NC_NOWRITE,
            MPI_COMM_WORLD, MPI_INFO_NULL, &ncid);

nc_inq_varid(ncid, "data", &varid);

nc_get_vara_int(ncid, varid, start, count, data);

nc_close(ncid);
```

You can use other calls to inquire about the dimensions too

- Facets
 - **Integration:** Easy to integrate with existing codes
 - **Visibility:** Not particularly transparent, can be hard to get good performance
 - **Flexible:** There is plenty of documented functionality
 - **Portable and supported:** Installed on most machines, relatively easy to compile and actively developed
 - **Extensible:** Open source but not designed to be user extensible as such.
 - **Documented:** Plenty of documentation, examples and tutorials available
- Disadvantages
 - Can be hard to get good parallel performance, whilst the framework promotes abstraction you can often end up working with lower level settings

- In this lecture we have considered the design of frameworks
- Some considerations to bear in mind when designing or selecting a framework
- Many of the suggestions apply equally well to writing good HPC code, even if it is not part of a framework.
- In the next lecture we will consider some frameworks which abstract parallelism and a more concrete look at the implementation of frameworks