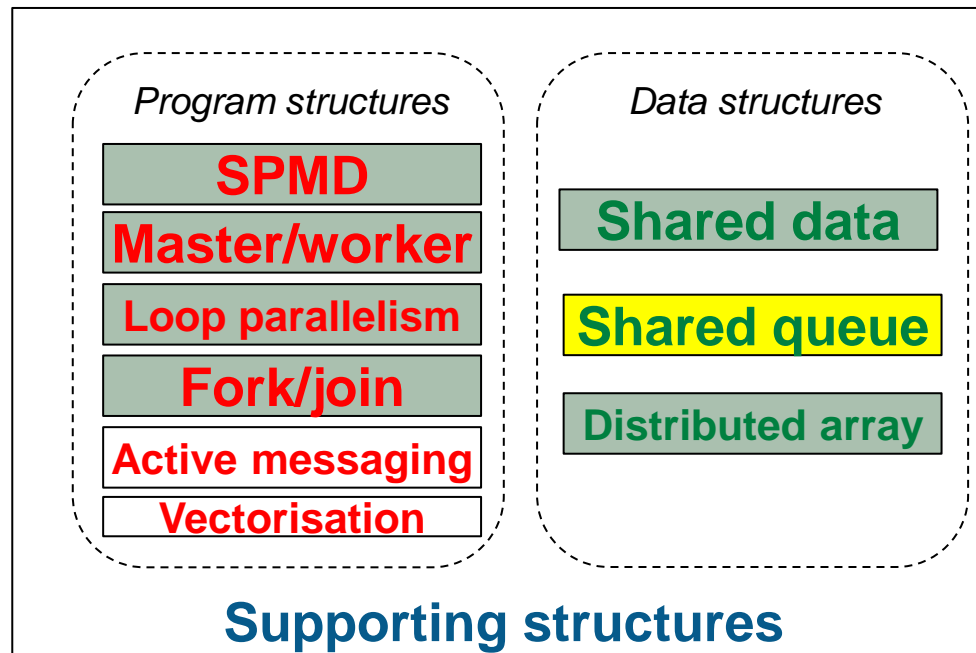


Parallel Design Patterns-L12

Shared Queue


Course Organiser: Dr Nick Brown
nick.brown@ed.ac.uk
Bayes Room 2.13



- Shared Queue is an *Implementation Pattern*
 - Based on a data type
- *The Problem:* How can concurrently-executing UEs safely share a queue data structure?



- Effective implementation of many parallel algorithms requires a queue that is to be shared among UEs
- An example we've already talked about is the “bag of tasks” in the Master-Worker pattern

- The queue is a FIFO data type





- Most of the important forces relate to the choice of **concurrency-control protocol**:

- One-at-a-time execution
- Non-interfering sets of operations
- Readers/Writers
- Splitting or Shrinking the Critical Section
- Nested Locks
- Application specific semantic relaxation

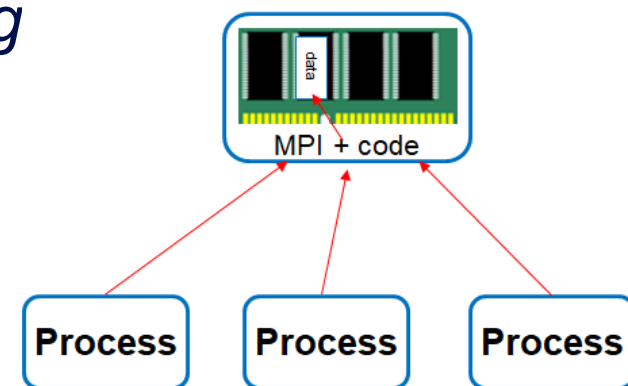
Simple but slow



Complex but fast


- Simple concurrency-control protocols provide greater clarity of abstraction and make it easier for the programmer to verify that the shared queue has been correctly implemented
 - *Aim for clarity first, then optimise*
- Concurrency-control protocols that encompass too much of the shared queue in a single synchronisation construct increase the chances UEs will remain blocked waiting to access the queue and will limit concurrency
- A concurrency-control protocol finely tuned to the queue and how it will be used increases the available concurrency, at the cost of more complicated, more error-prone synchronisation constructs

- Ideally the shared queue would be implemented as part of the target programming language
 - e.g. Java has an implementation available in `java.util.concurrent`
- No provided mechanism in common HPC languages (MPI, OpenMP)
- Most common use of shared queue is with *shared memory*
- Can be implemented in *message passing* by having the queue owned by one process, and putting and taking from the queue implemented by sending messages to and from the owner process



Apply the shared data pattern

- Define the ADT
- Choose the concurrency protocol



- The operations:
 - Put (*enqueue*)
 - Take (*dequeue*)
 - Other operations are possible, e.g. peek, takeall, clear, isEmpty
- Details:
 - What do you do when a queue is empty?
 - Block and wait for something to arrive
 - Could be used in Master-Worker with poison pill approach
 - Non-blocking queue: Return null or special value

- Implementing a shared queue can be tricky
 - but well-written, it can be re-used widely
- Choice of protocols
 - One-at-a-time execution
 - Non-interfering sets of operations
 - Readers/Writers
 - Splitting or Shrinking the Critical Section
 - Nested Locks
 - Application specific semantic relaxation

One at a time: Non-blocking

```
public class SharedQueue1 {
    class Node { //inner class defines list nodes {
        Object task;
        Node next;
        Node(Object task) {this.task = task; next = null;}
    }
    private Node head = new Node(null); //dummy node
    private Node last = head;

    public synchronized void put(Object task) {
        assert task != null: "Cannot insert null task";
        Node p = new Node(task);
        last.next = p;
        last = p;
    }

    public synchronized Object take() {
        //returns first task in queue or null if queue is empty
        Object task = null;
        if (!isEmpty()) {
            Node first = head.next;
            task = first.task;
            first.task = null;
            head = first;
        }
        return task;
    }
    private boolean isEmpty(){return head.next == null;} }
```

- A simple queue of ints, for illustration purposes:

```
void put (int i){  
    #pragma omp critical  
    ...  
    #pragma omp end critical  
}  
  
int take(){  
    #pragma omp critical  
    ...  
    #pragma omp end critical  
}
```


One at a time: Block on queue empty

```
public class SharedQueue2 {
    class Node {
        Object task;
        Node next;
        Node(Object task) {this.task = task; next = null;}
    }
    private Node head = new Node(null);
    private Node last = head;

    public synchronized void put(Object task) {
        assert task != null: "Cannot insert null task";
        Node p = new Node(task);
        last.next = p;
        last = p;
        notifyAll();
    }

    public synchronized Object take() {
        //returns first task in queue, waits if queue is empty
        Object task = null;
        while (isEmpty()) {
            try{wait();}catch(InterruptedException ignore){}
        }
        Node first = head.next;
        task = first.task;
        first.task = null;
        head = first;
        return task; } }
```

- Wait will release lock
 - Waits until notified
- notifyAll wakes all threads
 - In tern as lock on take method
- Pthreads has condition variables
 - Wait and signal

```
public class SharedQueue1 {
    class Node { //inner class defines list nodes {
        Object task;
        Node next;
        Node(Object task) {this.task = task; next = null;}
    }
    private Node head = new Node(null); //dummy node
    private Node last = head;

    public synchronized void put(Object task) {
        assert task != null: "Cannot insert null task";
        Node p = new Node(task);
        last.next = p;
        last = p;
    }

    public synchronized Object take() {
        //returns first task in queue or null if queue is empty
        Object task = null;
        if (!isEmpty()) {
            Node first = head.next;
            task = first.task;
            first.task = null;
            head = first;
        }
        return task;
    }
    private boolean isEmpty(){return head.next == null;} }
```

Non-interfering operations

```
public class SharedQueue3 {
    class Node {
        Object task;
        Node next;
        Node(Object task) {this.task = task; next = null;}
    }

    private Node head = new Node(null);
    private Node last = head;

    private Object putLock = new Object();
    private Object takeLock = new Object();

    public void put(Object task) {
        synchronized(putLock) {
            assert task != null: "Cannot insert null task";
            Node p = new Node(task);
            last.next = p; last = p;
        }
    }

    public Object take() {
        Object task = null;
        synchronized(takeLock) {
            if (!isEmpty()) {
                Node first = head.next;
                task = first.task;
                first.task = null;
                head = first;
            }
        }
        return task; } }
```

- Put and take are independent as do not access the same variables
- Therefore use different locks
- Only works for non blocking
- Could be two different mutexes in pthreads

- A simple queue of ints, for illustration purposes:

```
void put (int i){  
    #pragma omp critical(put)  
    ...  
    #pragma omp end critical(put)  
}  
  
int take(){  
    #pragma omp critical (take)  
    ...  
    #pragma omp end critical (take)  
}
```


Nested locks

```
public class SharedQueue4 {
    class Node {
        Object task; Node next;
        Node(Object task) {
            this.task = task; next = null;}
    }
    private Node head = new Node(null);
    private Node last = head;
    private int w;
    private Object putLock = new Object();
    private Object takeLock = new Object();

    public void put(Object task) {
        synchronized(putLock) {
            assert task != null: "Cannot insert null task";
            Node p = new Node(task);
            last.next = p; last = p;
            if(w>0) putLock.notify();
        }
    }
    public Object take() {
        Object task = null;
        synchronized(takeLock) {
            //returns first task in queue, waits if queue is empty
            while (isEmpty()) {
                try { synchronized(putLock){ w++; putLock.wait();w--; }
                } catch (InterruptedException error){assert false;}
            }
            Node first = head.next;
            task = first.task;
            first.task = null; head = first;
        }
        return task; } }
}
```

- Blocking on empty
- Waits on the putLock lock
- Need to be very careful to avoid deadlock

```
private Node last = head;
```

```
Rwlock rw_lock=new Rwlock();
```

```
public void put(Object task) {
```

```
    rw_lock.writeLock();  
    assert task != null: "Cannot insert null task";  
    Node p = new Node(task);  
    last.next = p; last = p;  
    rw_lock.release();  
}
```

```
public Object viewlast() {
```

```
    Object task = null;  
    rw_lock.readLock();  
    if (!isEmpty()) {  
        task=last.task;  
    }  
    rw_lock.release();  
    return task; } }
```

- Here *last* is used in both the functions
 - But one writes whilst the other reads
 - The reader can operate concurrently
 - Only one writer exclusively
- An example of this is rwlocks in pthreads

Shrinking the critical section

```
private Node last = head;

Rwlock rw_lock=new Rwlock();

public void put(Object task) {
    assert task != null: "Cannot insert null task";
    Node p = new Node(task);
    rw_lock.writeLock();
    last.next = p; last = p;
    rw_lock.release();
}

public Object viewlast() {
    Object task = null;
    rw_lock.readLock();
    if (!isEmpty()) {
        task=last.task;
    }
    rw_lock.release();
    return task; } }
```

- One central queue can be a bottleneck
 - Can we split this up so there is a queue per UE and distribute the contents?
- If my local queue becomes empty then a *take* might “steal” an element from a neighbour’s queue
- If my local queue becomes full then a *put* might add the element to a neighbour’s queue
- E.g. Allocating tasks to each UE to execute, queue these up and then allow for work stealing once completed.

- Shared Data:
 - Shared Queue pattern is an instance of Shared Data pattern.
- Master/Worker:
 - Shared Queue pattern is often used to represent the task queues in algorithms that use the Master/Worker pattern.
- Fork/Join pattern:
 - Thread-pool-based implementation of Fork/Join pattern is supported by this pattern such as the tasks of OpenMP we saw in the practical

- Idea: A shared queue encapsulates the synchronisation required inside an abstract data type
- Examples have been OO, but you can “encapsulate” inside put and take routines
- Different implementations can vary in performance and complexity
- Shared queue is a key component of various other parallel patterns