

Parallel Design Patterns-L14

Memory bound codes

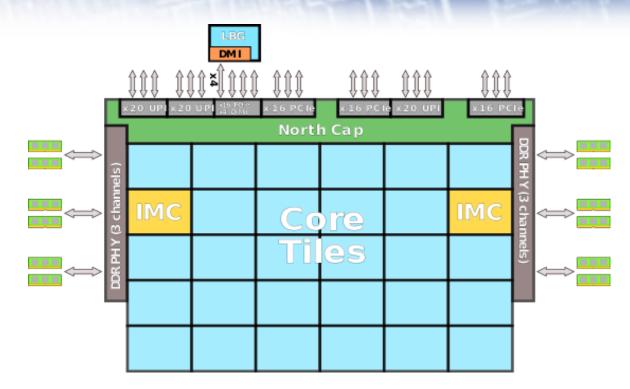
Spatial computing

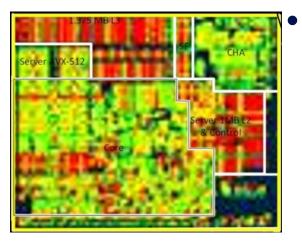
Course Organiser: Dr Nick Brown nick.brown@ed.ac.uk
Bayes room 2.13

The challenge with modern CPUs



- We have lots of compute
 - Many cores
 - Each core has significant FP performance



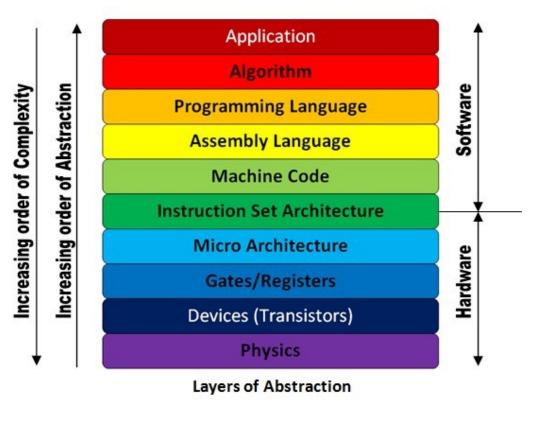


- Crucially the connection to DRAM is much more limited
 - So there is a very large cost in going "off-chip"
 - Caches can help, but there can be overheads with irregular memory access patterns
 - Many HPC codes are thus memory bound

Modern CPU



- CPUs expose themselves via an instruction set architecture (ISA)
 - E.g. x86, x86-64, A32, A64, SPARC, POWER
- A micro-architecture implements the ISA, is largely hidden and can change significantly between generations
 - Useful for the compiler to be aware of this and potentially more advanced programmers

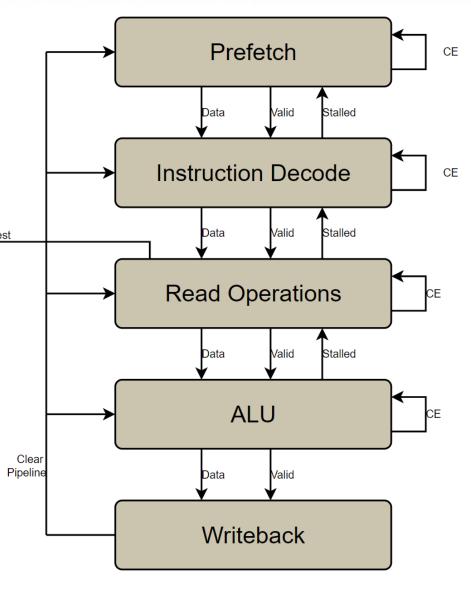


Traditional micro-architecture assumption



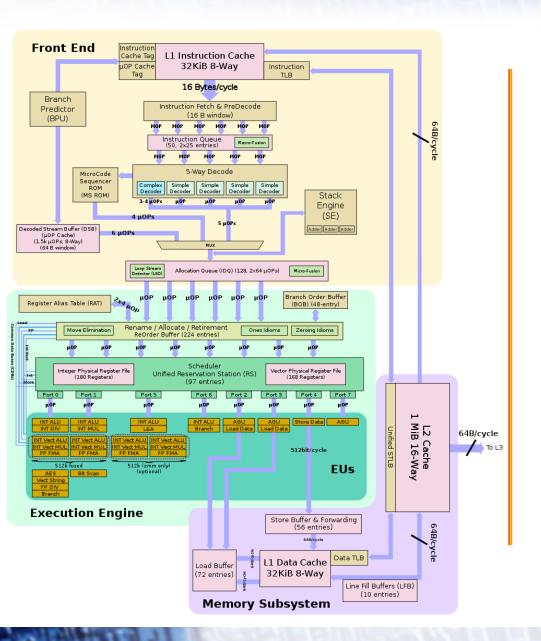
Some sort of pipelined
 approach where
 instructions start at the top
 and are executed a state at
 a time
 External halt request

- Assuming one cycle per stage, once pipelined is filled then one instruction retires (i.e. completes) per cycle
- You can see how things like cache misses can cause performance issues here



Out of Order execution





ld	Instruction	Issued	Executed
1	ld F6, 34(R2)	Υ	Υ
2	ld F2, 45(R3)	Υ	
3	multd F0, F2, F4	Y	
4	subd F8, F6, F2	Υ	
5	divd F10, F0, F6	Y	
6	addd F6, F8, F2	Υ	

Reorder buffer

	EU	Operation	OP1	OP2	
	Load1	ld	45(R3)		
*		multd	(RoB2)	(F4)	
*		subd	(RoB1)	(RoB2)	
*		divd	(RoB3)	(RoB1)	
*		addd	(RoB4)	(RoB2)	

Reservation station

Reorder Buffer (ROB)

Retirement Stage

Executed micro-ops, in order

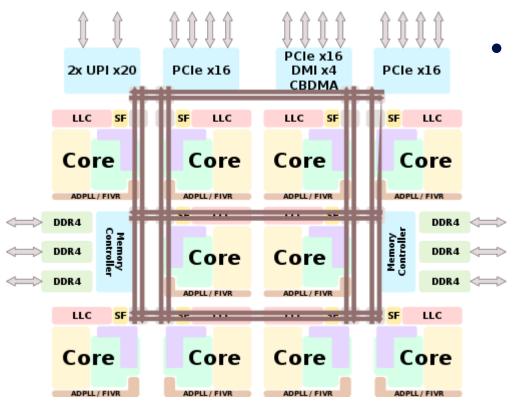
Reservation Station (RS)

Out-Of-Order (OOO)

As programmers, why care?



- The CPU really tries to help you out here
 - CPU caches
 - HW prefetchers that will look ahead and fetch the next data
 - Out of order execution to minimise the overhead of memory access



Therefore missing the cache doesn't always matter as the execution engine might be able to
 execute other instructions

- What gets us are delinquent misses, which stall the engine
- Sometimes limited what we can do about this

Demonstrating this in practice:



- With our practical one serial code (1024 by 1024 to 1e-3)
 - Compiled with GCC, -O3 and run on Skylake-X CPU

Code as provided

```
for(i=0; i<N; i++) {
  for (j=0; j<N; j++) {
    u_k[i][j]=u_kp1[i][j].....
}</pre>
```

Execution time	15.54 seconds
Number cycles	56.88 billion
Average IPC	2.45
MFLOP/s	5786.58
FPU occupancy	129.32%
Read from DRAM	125999.18 MB
Cache miss reads from DRAM	1.97 MB
% memory ops hit in L1	99.31%
% memory ops hit in L2	0.53%
% memory ops hit in L3	0.03%
% memory ops missed L3	0%
% cycles stalled due to memory	5.98%
Prefetcher hit in the L2 cache	51.39%

Loop order reversed (bad for cache)

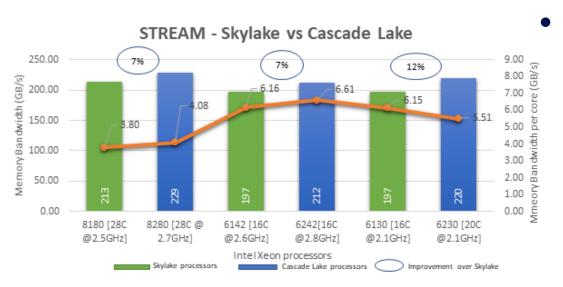
```
for(j=0; j<N; j++) {
  for (i=0; i<N; i++) {
    u_k[i][j]=u_kp1[i][j].....
}</pre>
```

Execution time	87.72 seconds
Number cycles	322.11 billion
Average IPC	0.59
MFLOP/s	1028.26
FPU occupancy	28%
Read from DRAM	127917.93 MB
Cache miss reads from DRAM	58.35 MB
% memory ops hit in L1	70.5%
% memory ops hit in L2	26.35%
% memory ops hit in L3	1.22%
% memory ops missed L3	2%
% cycles stalled due to memory	58.86%
Prefetcher hit in the L2 cache	29.06%

So what can we do about this?



- The previous example was an extreme one
 - And in this case improving our memory access will significantly help here.
 - But often we can't change our memory access pattern because it is what it is!

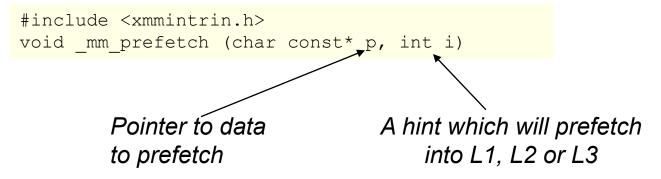


- CPU generations and technologies might help
 - But can be heavily dependent on HW improvements, and still leaving performance on the table
- GPUs might help, but still the same fundamental problem
 - Can sometimes hide memory access cost due to swapping warps and faster DDR memory is common-place. But doesn't always help!

Prefetching



 Can compliment or override the HW prefetcher by using SW prefetching via the _mm_prefetch function in your code



- Some compilers can do this automatically for you
 - I.e. With Intel compiler –opt-prefetch=n, where n is between 1 and 5 (higher being more prefetching)
- E.g. Integer Sort from NASA's Parallel Benchmark Suite

```
for( i=0; i<NUM_KEYS; i++ ) {
  key = key_array[i];
  key_buff2[bucket_ptrs[key >> shift]++] = key;
}
```

Access to key_array is predictable, but access to bucket_ptrs and key_buff2 is not

Modifying the code to use SW prefetching



```
#include <xmmintrin.h>
for( i=0; i<NUM KEYS+pf1+pf2; i++ ) {</pre>
  if (i < NUM KEYS) {
    key = key array[i];
    idx=i%tot distance;
                                         Prefetch into L1 cache
    keys buffer[idx]=key;
    bucket idx=key >> shift;
    idx buffer[idx]=bucket idx;
    mm prefetch(&bucket ptrs[bucket idx], MM HINT T0);
                               Work on i-pf1th key
  idx=(i-pf1);
  if (idx >= 0 \&\& idx < NUM KEYS) {
    idx=idx%tot distance;
    bucket idx=idx buffer[idx];
    value=bucket ptrs[bucket idx];
    buffer access[idx]=value++;
    bucket ptrs[bucket idx]=value;
    mm prefetch(&key buff2[value], MM HINT T0);
                                Work on i-(pf1+pf2)<sup>th</sup> key
  idx=(i-(pf1+pf2));
  if (idx >= 0) {
    idx=idx % tot distance;
    key buff2[buffer access[idx]]=keys buffer[idx];
```

Access to key_array is predictable, so that should be HW prefetched. Then do a SW prefetch for bucket_ptrs

pf1 iterations later, use the prefetched value from bucket_ptrs, increment and prefetch key buff2

pf2 iterations later use the prefetched value from key_buff2 and assign value to this

So what should the prefetch distance be?



Configuration	Runtime (s)
No prefetching	30.92
Intel compiler flag (n=3)	16.63
Intel compiler flag (n=5)	30.94
Prefetch distance 1	16.63
Prefetch distance 2	14.68
Prefetch distance 4	14.54
Prefetch distance 8	14.55
Prefetch distance 16	14.38
Prefetch distance 24	20.21

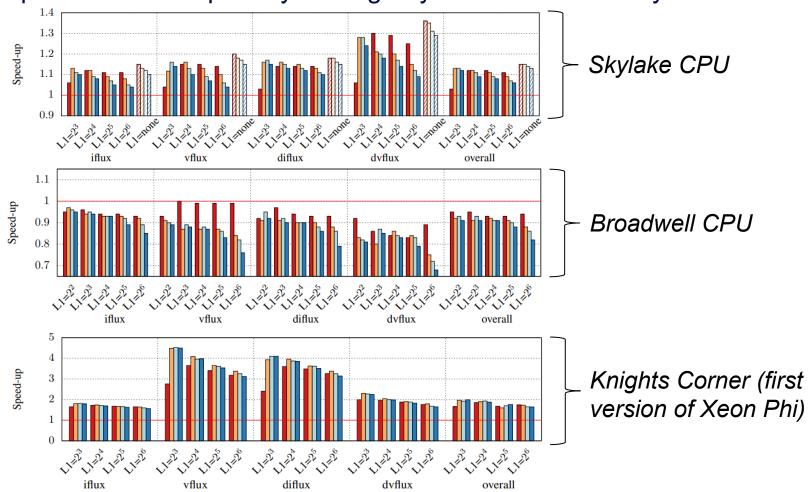
Intel compiler 2018, compiled with –O3 and run on Cirrus

- Compiler based
 prefetching is reasonable
 at the correct setting
- But can shave another 2 seconds off with our manual approach
 - But PF distance here is key
 - Adds considerable code complexity and might loose sequential equivalence
- Worth trying, but your results may vary!
 - More overall instructions need to be issued by the CPU
 - Might result in more work for the memory controller

Real-world code example



- Prefetching with an unstructured mesh code
 - A paper "Software prefetching for unstructured mesh applications" published a couple of years ago by Warwick University



Advice for us as programmers



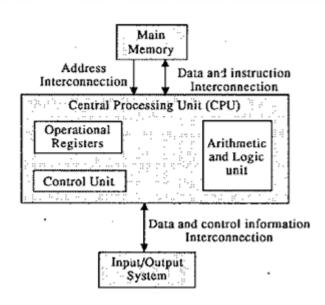
- Profile your codes to understand if memory bound
- If they are memory bound then could try a number of things
 - Experiment with different memory layouts (e.g. AoS or SoA)
 - If you have lots of indirect memory accesses then experiment with SW prefetching, first compiler based and then possibly manual prefetching if necessary
 - Distributing to another NUMA region (e.g. another CPU in the same memory space) or across a distributed machine might help
 - Effectively increasing the number of memory controllers and channels you have in the global system

 Or wait until a CPU comes along with more memory bandwidth/lower latency!

But there is (potentially) another way!



- Fundamentally we have lots of transistors in a chip and want these to all be doing something useful all of the time
 - Put quite simply, the Von Neumann architecture isn't how the electronics works
 - Current CPUs (and GPUs) give the illusion of Von Neumann but implement this more efficiently



– Where these two views of the world meet can often be a pain point!

- But there is another way! What if we unify the view exposed to the programmer and how the electronics operates?
 - By designing this around the flow of data we can avoid the memory bound nature of our code?

Step in: Reconfigurable (spatial) computing



- Creating a custom chip, an ASIC, costs millions of dollars
 - So not realistic to create a chip per application!



- Instead we can use a Field Programmable Gate Array (FPGA) which can be programmed at the electronic level to behave pretty much like anything we want
 - Such as the electronic implementation of our application





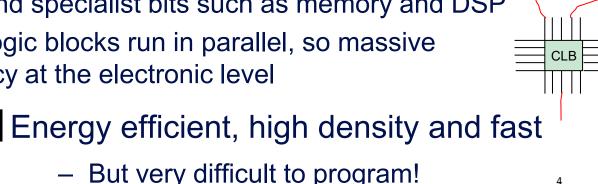
What is an FPGA

CLB

CLB

Data values are

- A semiconductor based around a matrix of configurable logic blocks (CLBs) which are themselves connected via programmable interconnects.
 - Allow these blocks to be wired together in many different ways to form lots of different circuits and logic gates.
 - First invented by Xilinx in 1985
 - FPGAs contain lots of interconnect so you can do lots of wiring and specialist bits such as memory and DSP
 - All these logic blocks run in parallel, so massive concurrency at the electronic level

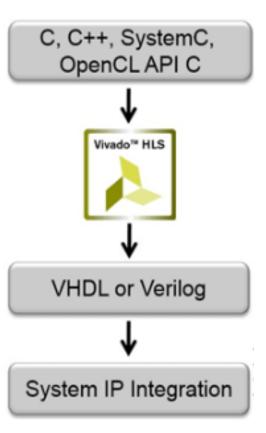




- But very difficult to program!
- Becoming available in public clouds



Was difficult, now is different:



 Program directly using a hardware description language such as VHDL or Verilog (the machine code level)

2. Write code in C or C++ and use high level synthesis (HLS) to synthesise this down to VHDL or Verilog

3. Combine point 2 with OpenCL on the host and automatic integration with the chip

A simple example via high level synthesis



```
void sum_kernel(float * input, float * result, float add_val, int num_its) {
    #pragma HLS INTERFACE m_axi port=input offset=slave
    #pragma HLS INTERFACE m_axi port=result offset=slave
    #pragma HLS INTERFACE s_axilite port=add_val bundle=control
    #pragma HLS INTERFACE s_axilite port=num_its bundle=control
    #pragma HLS INTERFACE s_axilite port=return bundle=control

float sum=0;
for (unsigned int i=0;i<num_its;i++) {
    float d=input[i] + add_val;
    sum+=d;
}
*result=sum;
}</pre>
```

```
> v++ -t hw --config design.cfg -03 -c -k sum_kernel -o'sum.hw.xo' device.cpp
```

- Compiles this code into the underlying hardware description language
 - Over 20,000 lines of VHDL!
 - On disk these are .xo files which can be thought of similarly to object files

	·				
Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	57	-
FIFO	-	-	-	-	-
Instance	4	2	1454	1526	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	220	-
Register	0	-	315	32	-
Total	4	2	1769	1835	0
Available	4032	9024	2607360	1303680	960
Available SLR	1344	3008	869120	434560	320
Utilization (%)	~0	~0	~0	~0	0
Utilization SLR (%)	~0	~0	~0	~0	0

Driving the FPGA from the host



```
#define DATA SIZE 1000000
std::vector<float, aligned allocator<float>> input data(DATA SIZE), result data(1);
cl::Context * context=new cl::Context(device, NULL, NULL, NULL);
// Create the command queue
cl::CommandQueue * command queue=new cl::CommandQueue(*context, device, CL QUEUE PROFILING ENABLE);
// Create a handle to the kernel
cl::Kernel * sum kernel=new cl::Kernel(*program, "sum kernel");
unsigned int base size = DATA SIZE * sizeof(float);
cl::Buffer * buffer input=new cl::Buffer(*context, CL MEM USE HOST PTR, base size, input data.data());
cl::Buffer * buffer result=new cl::Buffer(*context, CL MEM USE HOST PTR, sizeof(float), result data.data());
// Set kernel arguments
                                                                                  (1) Copy input data from CPU
                                                                                    onto FPGA via PCle
sum kernel->setArg(0, *buffer input);
sum kernel->setArg(1, *buffer result);
                                                                                                      FPGA
sum kernel->setArg(2, (float) 50.0);
sum kernel->setArg(3, DATA SIZE);
                                                                                 (3) Copy result data from FPGA to
                                                                                                     (2) Execute HLS kernels on FPGA
// Queue migration of memory objects from host to device
command queue->enqueueMigrateMemObjects({*buffer input, *buffer result}, 0);
// Queue kernel execution
command gueue->engueueTask(*sum kernel);
// Queue copy result data back from kernel
command queue->enqueueMigrateMemObjects({*buffer result}, CL MIGRATE MEM OBJECT HOST);
// Wait for queue to complete
command queue->finish();
```

```
> g++ -03 -std=c++11 -o 'host' 'host.cpp' -lpthread -lrt -lstdc++ -lOpenCL
> ./host
Elapsed time: 0.024823 secs, (0.000869 transfer on, 0.023954 execution, 0.000000 transfer off)
```

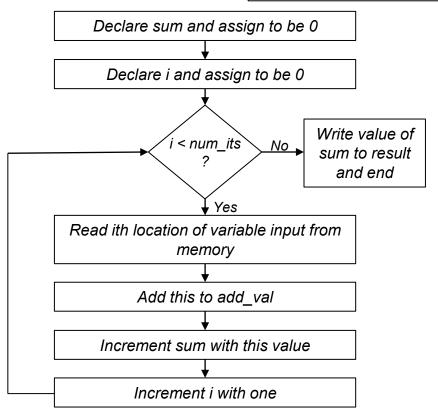
Where's the spatial computing?



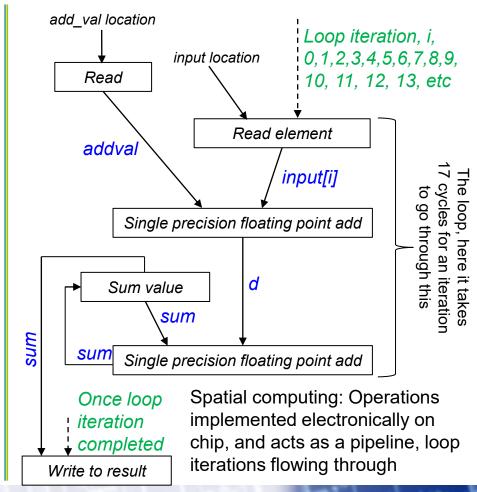
Temporal computing

```
float sum=0;
for (unsigned int i=0;i<num_its;i++) {
  float d=input[i] + add_val;
  sum+=d;
}
*result=sum;</pre>
```

Spatial computing (dataflow)

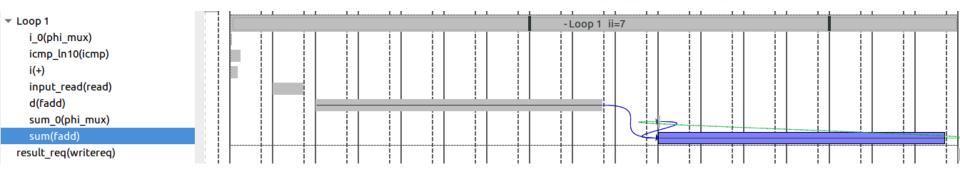


Temporal computing: Can be thought of like a flowchart, with the PE (e.g. CPU or GPU) executing one stage after another



This directly impacts our example...





float sum=0;

sum+=d;

for (int i=0;i<num_its;i++) {
 float d=input[i] + add val;</pre>

Pipeline depth of 17 cycles Initiation Interval of 7 cycles

There is a spatial dependency here

```
of 7 cycles
```

- Effectively a dependency between the first and last pipeline stage
 - The last state depends directly on the value generated by the first stage and feeds into this
- So we can only run a loop iteration every seven cycles
 - Instead want an iteration per cycle or else are throwing away lots of potential performance!
- CPU executes this in 0.004339 seconds (approx. 5 times faster!)

Removing the spatial dependency

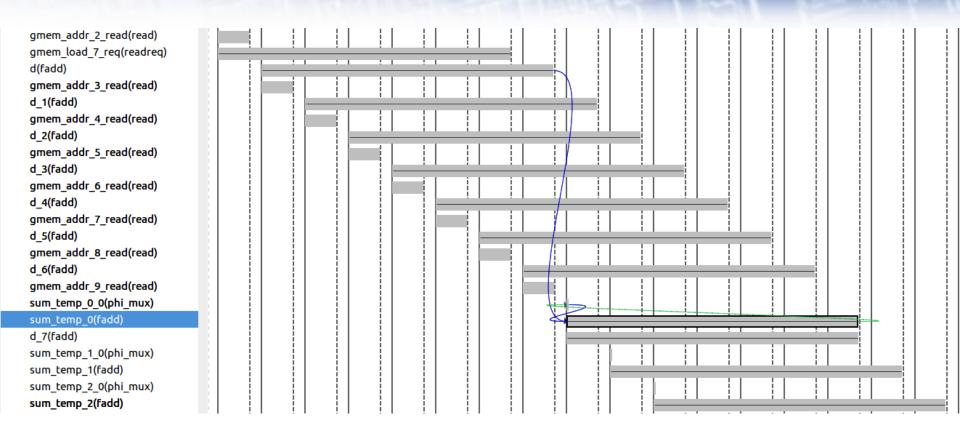


```
void sum kernel(float * input, float * result, float add val, int num its) {
    #pragma HLS INTERFACE m axi port=input
    #pragma HLS INTERFACE m axi port=result
    #pragma HLS INTERFACE s axilite port=add val bundle=control
    #pragma HLS INTERFACE s axilite port=num its bundle=control
    #pragma HLS INTERFACE s axilite port=return bundle=control
    float sum temp[8];
    #pragma HLS array partition variable=sum temp complete
    for (unsigned int i = 0; i < 8; i++) {
    #pragma HLS unroll
      sum temp[i]=0;
    const unsigned int base its=num its / 8, remainder its=num its - (base its*8);
    for(unsigned int i=0; i < base its; i++) {</pre>
    #pragma HLS PIPELINE II=8
      for (unsigned int j=0; j<8; j++) {
        float d=input[(i*8)+j] + add val;
        sum temp[j] += d;
    for (unsigned int i=0;i<remainder its;i++) {</pre>
    #pragma HLS unroll
      float d=input[i+(base its*8)] + add val;
      sum temp[i]+=d;
    for (unsigned int i = 0; i < 8; i++) {
    #pragma HLS unroll
      *result+=sum temp[i];
```

Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	1015	-
FIFO	-	-	-	-	-
Instance	2	4	1400	1433	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	1043	-
Register	0	-	2841	32	-
Total	2	4	4241	3523	0
Available	4032	9024	2607360	1303680	960
Available SLR	1344	3008	869120	434560	320
Utilization (%)	~0	~0	~0	~0	0
Utilization SLR (%)	~0	~0	~0	~0	0

How's this make a difference?





- Unrolling this inner loop and using a temporary variable to avoid the conflict on the 7 cycle floating point addition
 - Now runtime is 0.004722 seconds
 - Very slightly slower than the CPU

```
for(unsigned int i=0; i<base_its; i++) {
    #pragma HLS PIPELINE II=8
    for (unsigned int j=0; j<8; j++) {
        float d=input[i+j] + add_val;
        sum_temp[j]+=d;
    }
}</pre>
```

Removing read dependencies on "input"



```
struct packed data { float data[8]; };
void sum kernel(struct packed data * input, float * result, float add val, int num its) {
    #pragma HLS DATA PACK variable=input
    #pragma HLS INTERFACE m axi port=input
    #pragma HLS INTERFACE m axi port=result
    #pragma HLS INTERFACE s axilite port=add val bundle=control
    #pragma HLS INTERFACE s axilite port=num its bundle=control
    #pragma HLS INTERFACE s axilite port=return bundle=control
   float sum temp[8];
    #pragma HLS array partition variable=sum temp complete
   for (unsigned int i = 0; i < 8; i++) {
    #pragma HLS unroll
      sum temp[i]=0;
   const unsigned int base its=num its / 8, remainder its=num its - (base its*8);
   for(unsigned int i=0; i < base its; i++) {</pre>
    #pragma HLS PIPELINE II=8
      struct packed data v=input[i];
      for (unsigned int j=0; j<8; j++) {
        float d=v.data[j] + add val;
        sum temp[j]+=d;
    struct packed data v=input[base its];
   for (unsigned int i=0;i<remainder its;i++) {</pre>
    #pragma HLS unroll
      float d=v.data[i] + add val;
      sum temp[i]+=d;
    for (unsigned int i = 0; i < 8; i++) {
    #pragma HLS unroll
      *result+=sum temp[i];
```

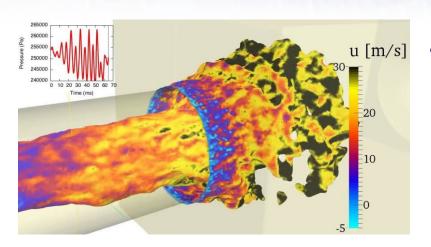
	• ' '				
Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	1246	-
FIFO	-	-	-	-	-
Instance	16	8	2891	2977	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	967	-
Register	0	-	2649	32	-
Total	16	8	5540	5222	0
Available	4032	9024	2607360	1303680	960
Available SLR	1344	3008	869120	434560	320
Utilization (%)	~0	~0	~0	~0	0
Utilization SLR (%)	1	~0	~0	1	0

Runtime is 0.002399 seconds

- 0.000829 transfer, and 0.001570 execution
- About twice as fast as the CPU and ten times faster than our initial version!

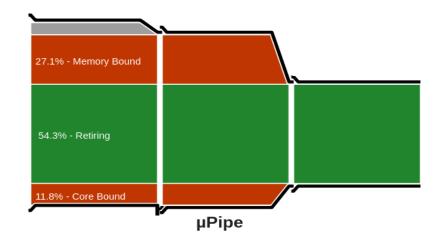
Alya example





 Developed by BSC for computational mechanics code used to solve complex coupled multi-physics, multi-scale, and multi-domain problems

- Nastin module models incompressible flow and is a key component of many run configurations.
- Building of the matrix, which is used to solve Navier-Stokes equations, represents 64% of the overall model runtime
 - Only 54% of cycles doing useful work,
 27% memory bound and 12% core bound



Routine	Time	FLOPs per element
calculate_transients	3.2%	0
calculate_cartesian_derivatives	5.4%	664
calculate_gauss_point_values	8.9%	400
calculate_tau_and_tim	2.1%	76
calculate_element_matricies	11%	416
calculate_convective_term_and_RHS	40%	3936
calculate_viscous_term	26%	1540
perform_assembly_in_global_system	3.4%	20

Optimising bottom up



```
for (int i=0; i<N; i++) {
    double d=x*y;
    double j=d*z;
    double p=d*j
    result=p;
}

for (int i=0; i<N; i++) {
    #pragma HLS pipeline II=1
    double d=x*y;
    double j=d*z;
    double p=d*j
    result=p;
}

Loop pipelining</pre>
```

```
for (int i=0; i<N; i++) {
#pragma HLS pipeline II=1
    double d=x*y;
    double j=d*z;
    double p=d*j
    result=p;
}</pre>
for (int i=0; i<N; i++) {
#pragma HLS pipeline II=1
#pragma UNROLL FACTOR=4
    double d=x*y;
    double j=d*z;
    double p=d*j
    result=p;
}
Loop unrolling
```

```
double val=0;
for (int i=0; i<N; i++) {
    #pragma HLS pipeline II=1
      val=val+external[i];
}</pre>
```

Spatial dependency

```
for (int i=0; i<N; i++) {
#pragma HLS pipeline II=1
    double d=external_data[i];
    double j=external_data[i+1];
    ...
}</pre>
```

Conflict on external port to HBM2/DDR memory

```
double local_data[M];
for (int i=0; i<N; i++) {
    #pragma HLS pipeline II=1
        local_data[i-2]=a;
        local_data[i-1]=b;
        double v=local_data[i];
}</pre>
```

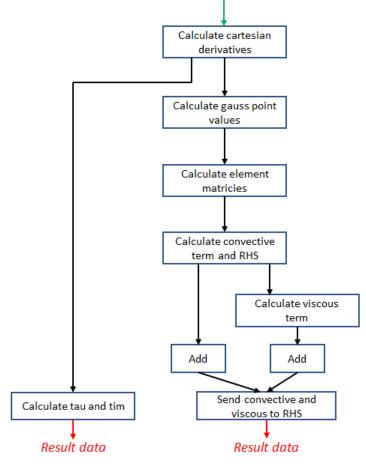
Conflict on (dual-ported) on-chip BRAM memory

Optimising top down too

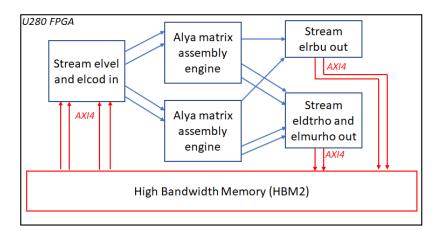


Description	Runtime (ms)	% of CPU
1 core Xeon CPU	351.05	-
24 cores Xeon CPU	61.72	-
Initial FPGA port	15714.99	0.39%
Optimised top down	1508.60	4.09%
Optimised bottom up	284.04	21.73%
Optimised external memory access	26.67	231.42%





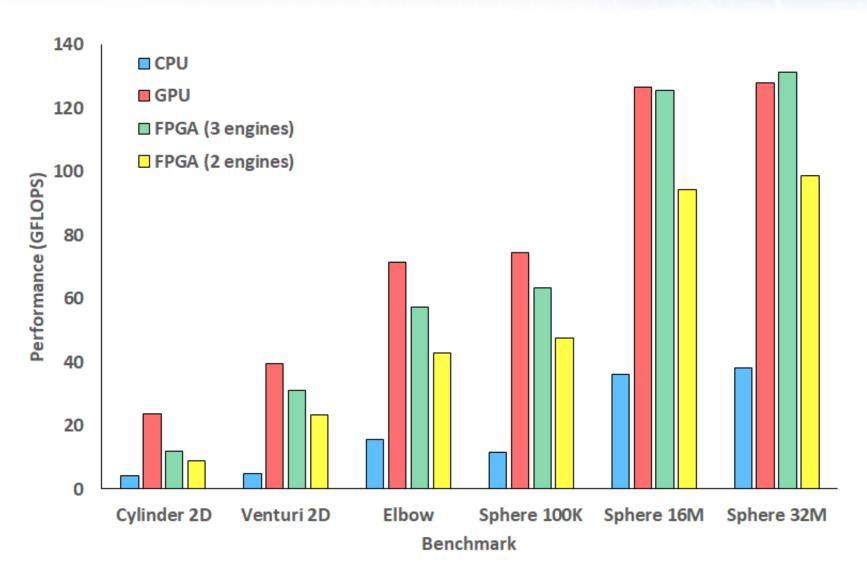
Input data



 Idea is to keep our compute fed with data, by running parts of the chip on compute and others on data loading/reordering

Performance against CPU and GPU

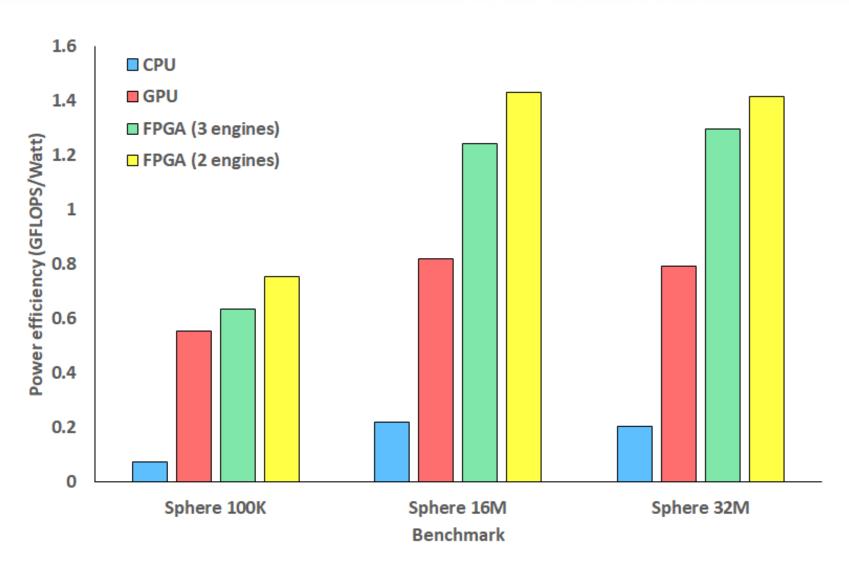




Xilinx Alveo U280 vs V100 GPU and Xeon Platinum (24 core) Cascade Lake CPU

Power efficiency against CPU and GPU



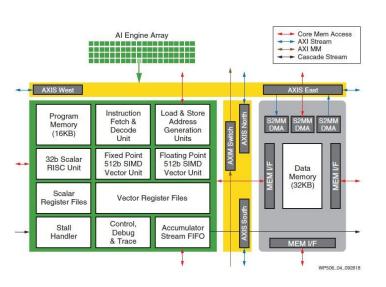


Xilinx Alveo U280 vs V100 GPU and Xeon Platinum (24 core) Cascade Lake CPU

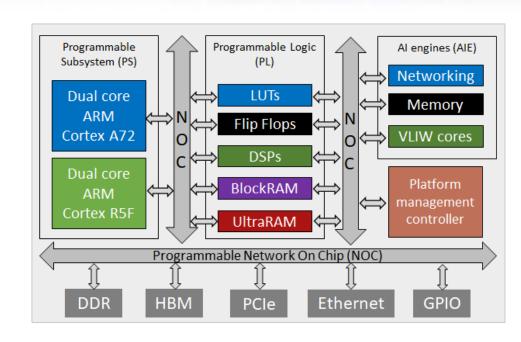
Adaptive Compute Acceleration Platform



- Combines reconfigurable fabric with standard CPUs and AI engines for vectorised acceleration
 - Al engines planned in AMD's next generation CPUs too



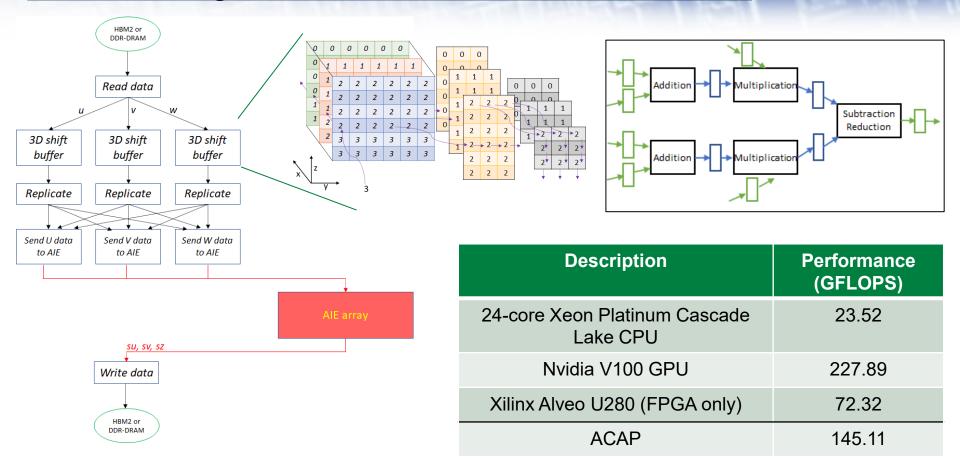
For single precision floating point, 3.8 TFLOPs theoretical peak performance total on the chip



- 400 AI engines running at 1.2GHz
 - 7-way VLIW cores
 - 16KB of program memory and 32KB data
 - Can access neighbouring core's memory
 - Two 32-bit input and two 32-bit output streams
 - Therefore can read 128-bit every 4 cycles
 - 384-bit cascade stream per core

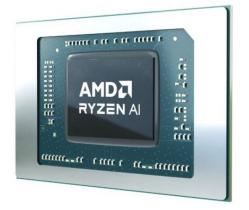
Accelerating advection on the AIEs



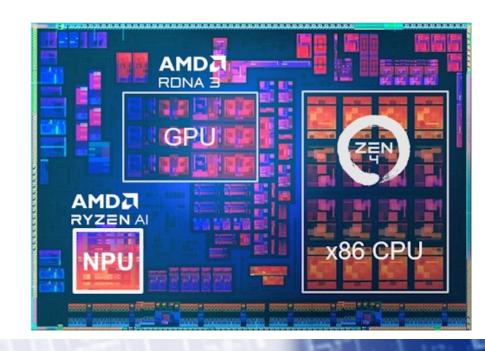


- Accelerating atmospheric advection on the ACAP
 - Idea is to use the reconfigurable fabric for bespoke memory accesses to specialise that aspect
 - Use AIEs for floating point arithmetic





- AMD have recently released a CPU with AI engines packaged in the chip
 - Currently aimed at laptop market, but potentially demonstrates direction of travel here
 - And have recently released Riallto which is a framework for programming this chip
- Focus is machine learning
 - Neural Processing Unit (NPU)
 but this is just a bunch of AIEs
 - Also released lots of compilation tooling too so that people can develop their own frameworks
 - Strong potential for HPC here



Spatial computing summary



- A very interesting approach and coming back in fashion
 - Future HPC architectures
 - Future embedded devices
 - Ways of bringing the computation to the data in challenging environments
 - Could especially help with memory bound codes
- But difficult to program as you also need to think spatially and ensure that the chip is kept as busy as possible
 - Requires thinking about problems in a data-flow manner
 - Compilation of code can be very time consuming
 - Active area of research to address these challenges, but currently there is nothing that solves the issues completely.
 - Jury is still out regarding their utility for HPC codes. Xilinx are currently selling them very heavily for machine learning applications