Practical two

Wrap up

Strong scaling

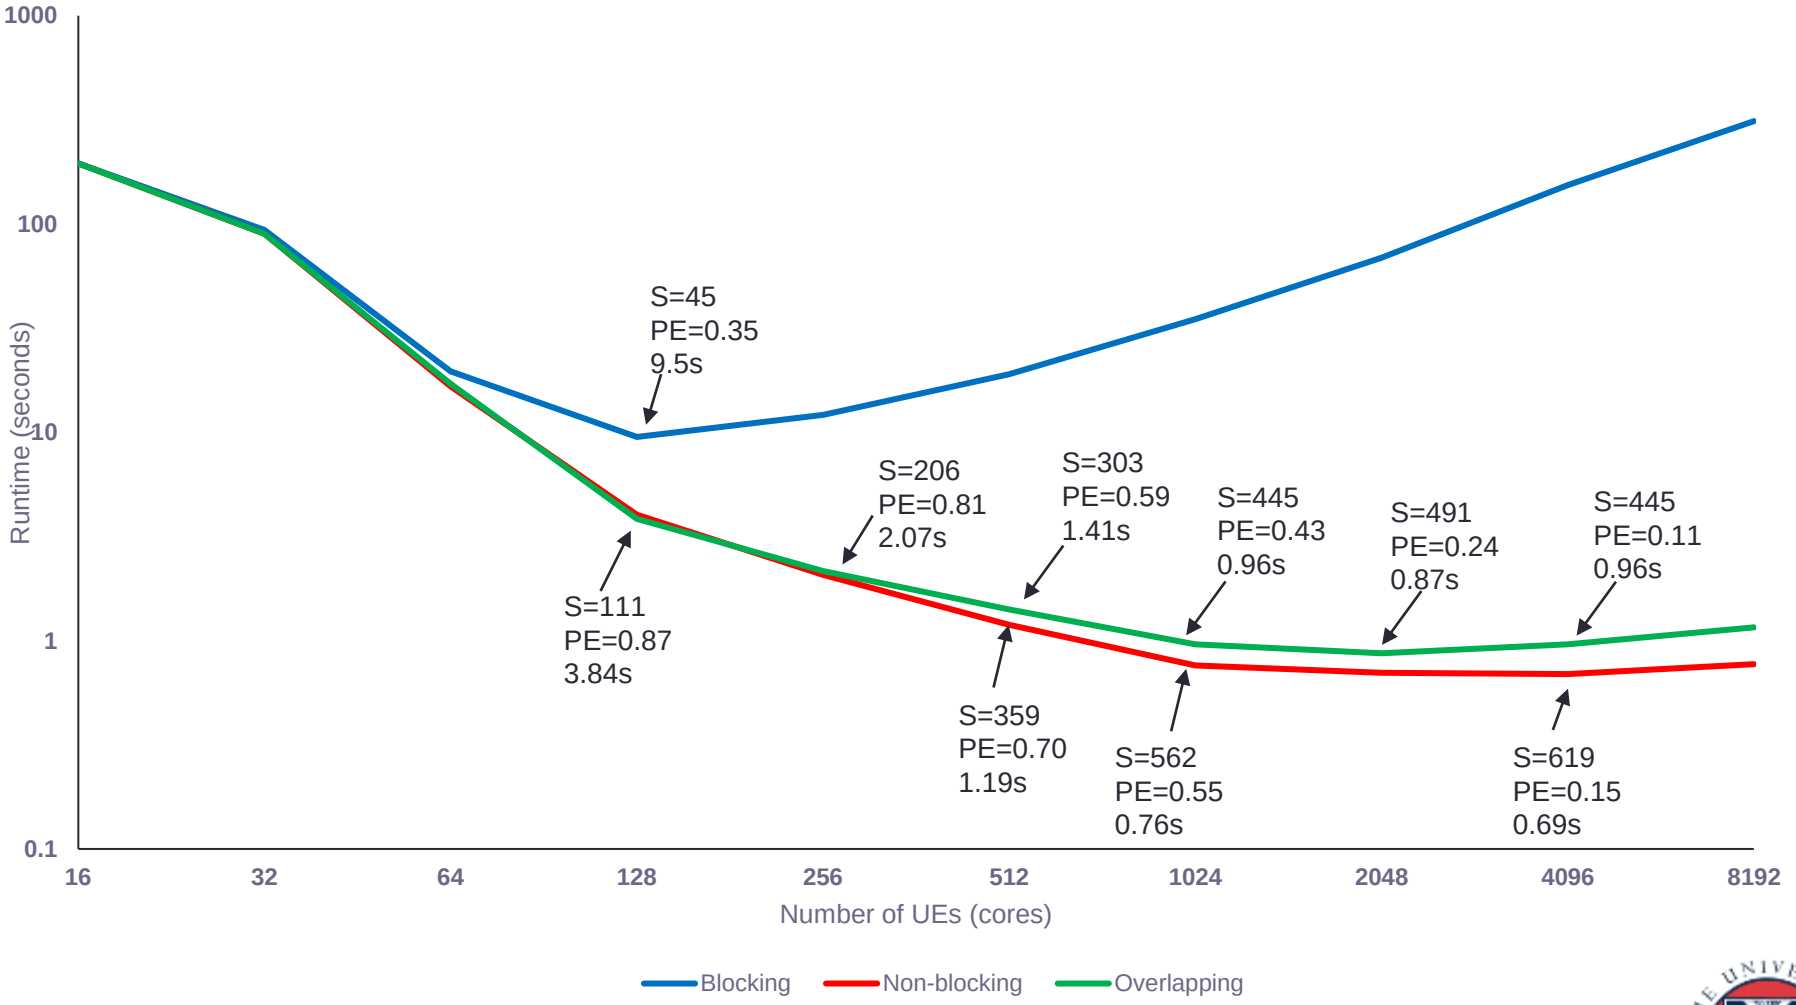

epcc

THE UNIVERSITY OF EDINBURGH

ARCHER2 vs ARCHER

Cores	ARCHER	ARCHER2
1	922.58	427.27
16	125.26	195.4
32	62.84	93.77
64	33.99	19.64
128	16.23	9.5
256	16.5	12.11

blocking

Cores	ARCHER	ARCHER2
1	922.58	427.27
16	124.63	195.26
32	61.01	89.94
64	29.62	16.61
128	8.82	4.02
256	5.22	2.07
512	4.45	1.19
1024	3.92	0.76

Non-blocking

epcc

THE UNIVERSITY OF EDINBURGH

ARCHER2 vs ARCHER

Cores	ARCHER	ARCHER2	ARCHER2 widely
1	922.58	427.27	-
16	125.26	195.4	34.91
32	62.84	93.77	29.11
64	33.99	19.64	13.82
128	16.23	9.5	13.37
256	16.5	12.11	-

blocking

Cores	ARCHER	ARCHER2	ARCHER2 widely
1	922.58	427.27	-
16	124.63	195.26	33.18
32	61.01	89.94	26.87
64	29.62	16.61	9.67
128	8.82	4.02	4.24
256	5.22	2.07	-
512	4.45	1.19	-
1024	3.92	0.76	-

Non-blocking

--distribution=block:block will fill up each NUMA region one by one

So will allocate to core 0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15 etc….

If we remove it will allocate widely

So will allocate to core 0,16,32,48,64,80,96,112,1,17,33,49,65,81,97,113 etc….

Therefore for small numbers of cores this is better as we use more NUMA regions and hence memory links (improve our memory bandwidth)

But when we use all the cores on the CPU then there is lots of cross NUMA communication (16 cores per NUMA region, so works badly here for us) and results in overhead. Also groups of 4 cores share an L3 cache.

Likely could get slightly higher performance with 32 and 64 cores if we pinned those in the queue submission file due to these reasons

epcc

THE UNIVERSITY OF EDINBURGH

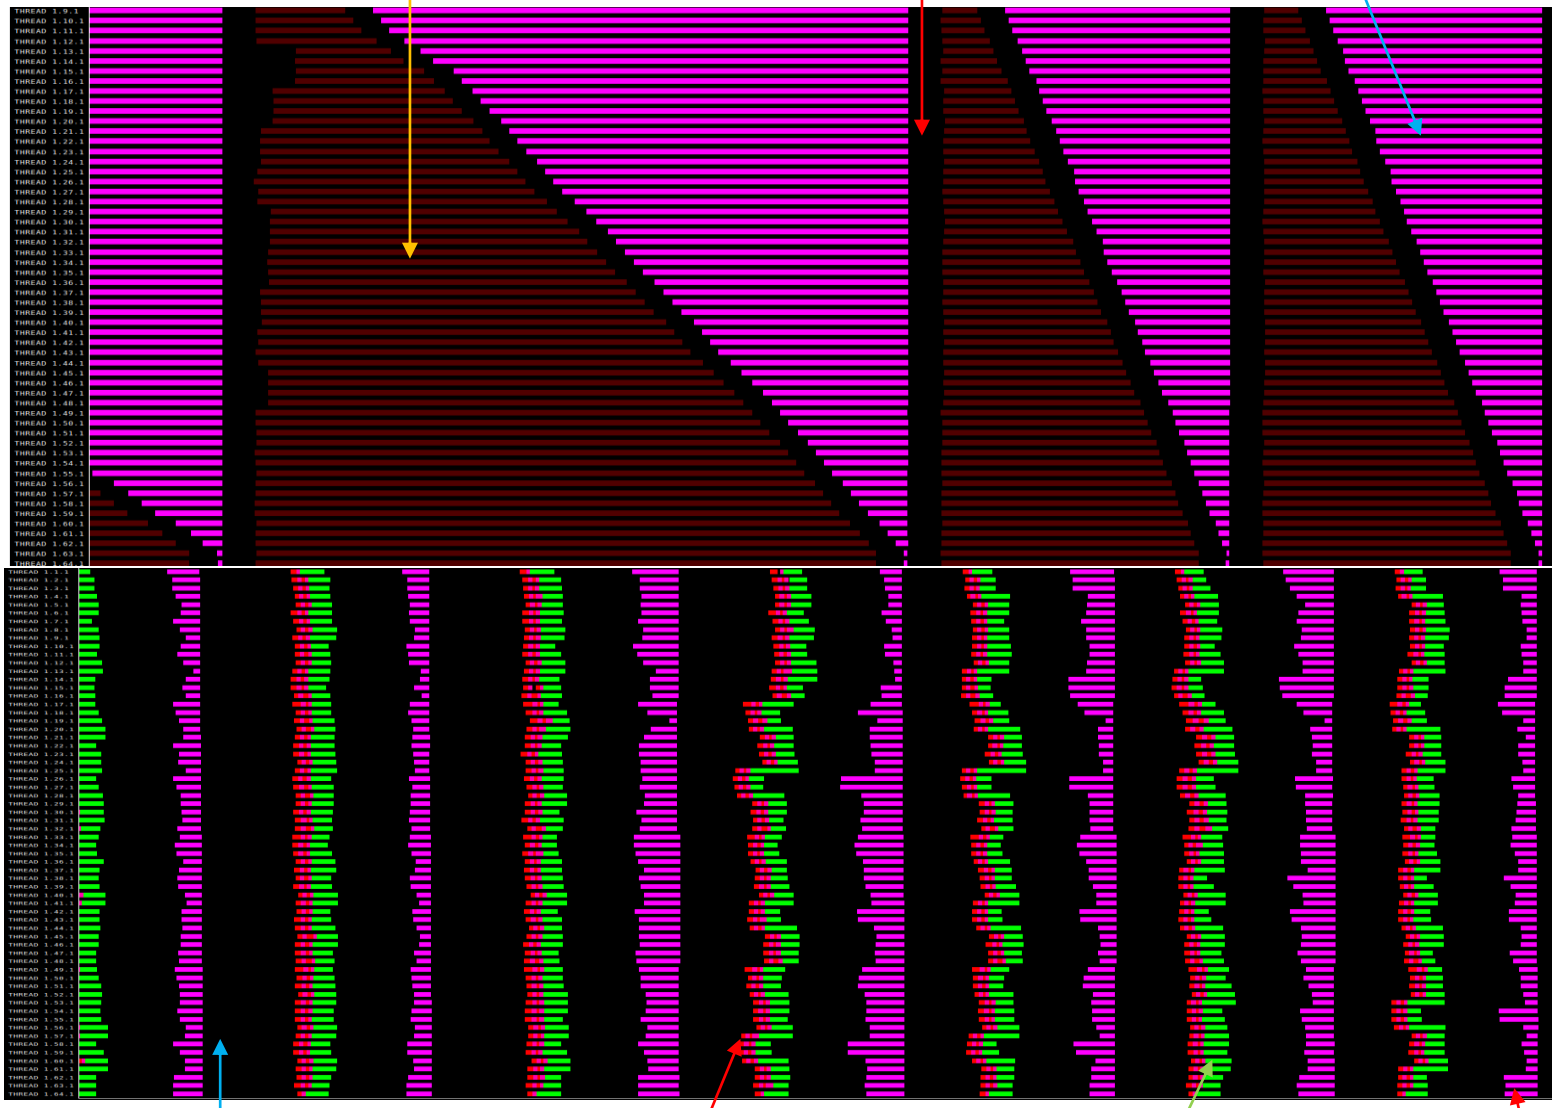
Sendrecv
Computation
Allreduce
Blocking parallel version
Non-blocking parallel version
Computation
Isend
Irecv
Waitall
Allreduce

	Outside MPI	MPI_Allreduce	MPI_Comm_rank	MPI_Comm_size	MPI_Init	MPI_Finalize	MPI_Sendrecv
THREAD 1.1.1	31.24 %	67.19 %	0.00 %	0.02 %	0.01 %	1.09 %	0.44 %
THREAD 1.2.1	19.09 %	66.65 %	0.00 %	0.02 %	12.79 %	0.43 %	1.02 %
THREAD 1.3.1	19.03 %	66.09 %	0.00 %	0.02 %	12.78 %	0.44 %	1.64 %
THREAD 1.4.1	19.03 %	65.50 %	0.00 %	0.01 %	12.78 %	0.44 %	2.23 %
THREAD 1.93.1	29.73 %	14.36 %	0.00 %	0.01 %	0.68 %	1.03 %	54.18 %
THREAD 1.94.1	29.89 %	13.80 %	0.00 %	0.02 %	0.48 %	1.04 %	54.77 %
THREAD 1.95.1	30.09 %	13.23 %	0.00 %	0.02 %	0.29 %	1.05 %	55.33 %
THREAD 1.96.1	30.19 %	13.23 %	0.00 %	0.02 %	0.10 %	1.05 %	55.41 %

Blocking parallel version

	Outside MPI	MPI_Isend	MPI_Irecv	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank	MPI_Comm_size	MPI_Init	MPI_Finalize
THREAD 1.1.1	66.08 %	0.42 %	0.39 %	1.41 %	29.49 %	0.00 %	0.05 %	0.01 %	2.16 %
THREAD 1.2.1	41.85 %	0.72 %	0.57 %	1.35 %	29.23 %	0.00 %	0.04 %	25.44 %	0.79 %
THREAD 1.3.1	41.87 %	0.73 %	0.57 %	1.28 %	29.26 %	0.00 %	0.05 %	25.43 %	0.80 %
THREAD 1.4.1	41.71 %	0.73 %	0.58 %	1.39 %	29.32 %	0.00 %	0.03 %	25.42 %	0.81 %
THREAD 1.93.1	62.93 %	0.70 %	0.56 %	1.24 %	31.20 %	0.00 %	0.03 %	1.33 %	2.00 %
THREAD 1.94.1	63.19 %	0.71 %	0.57 %	1.23 %	31.27 %	0.00 %	0.03 %	0.96 %	2.03 %
THREAD 1.95.1	63.57 %	0.71 %	0.56 %	1.18 %	31.30 %	0.00 %	0.04 %	0.59 %	2.04 %
THREAD 1.96.1	63.74 %	0.38 %	0.32 %	1.79 %	31.45 %	0.00 %	0.04 %	0.22 %	2.06 %

Non-blocking parallel version

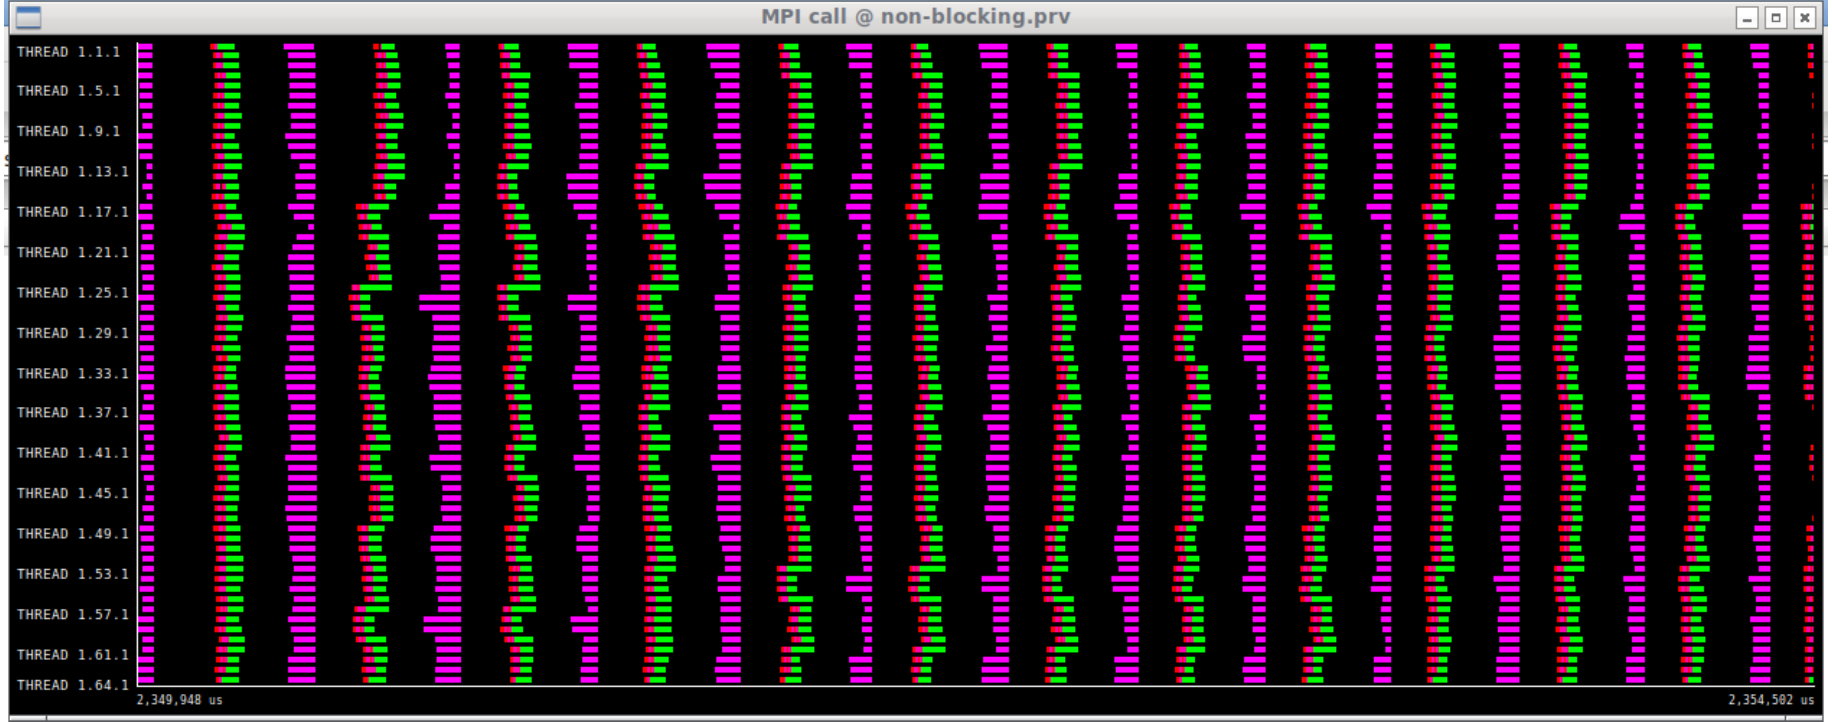
MPI call @ non-blocking.prv
THREAD 1.1.1
THREAD 1.5.1
THREAD 1.9.1
THREAD 1.13.1
THREAD 1.17.1
THREAD 1.21.1
THREAD 1.25.1
THREAD 1.29.1
THREAD 1.33.1
THREAD 1.37.1
THREAD 1.41.1
THREAD 1.45.1
THREAD 1.49.1
THREAD 1.53.1
THREAD 1.57.1
THREAD 1.61.1
THREAD 1.64.1
2,349,948 us
2,354,502 us

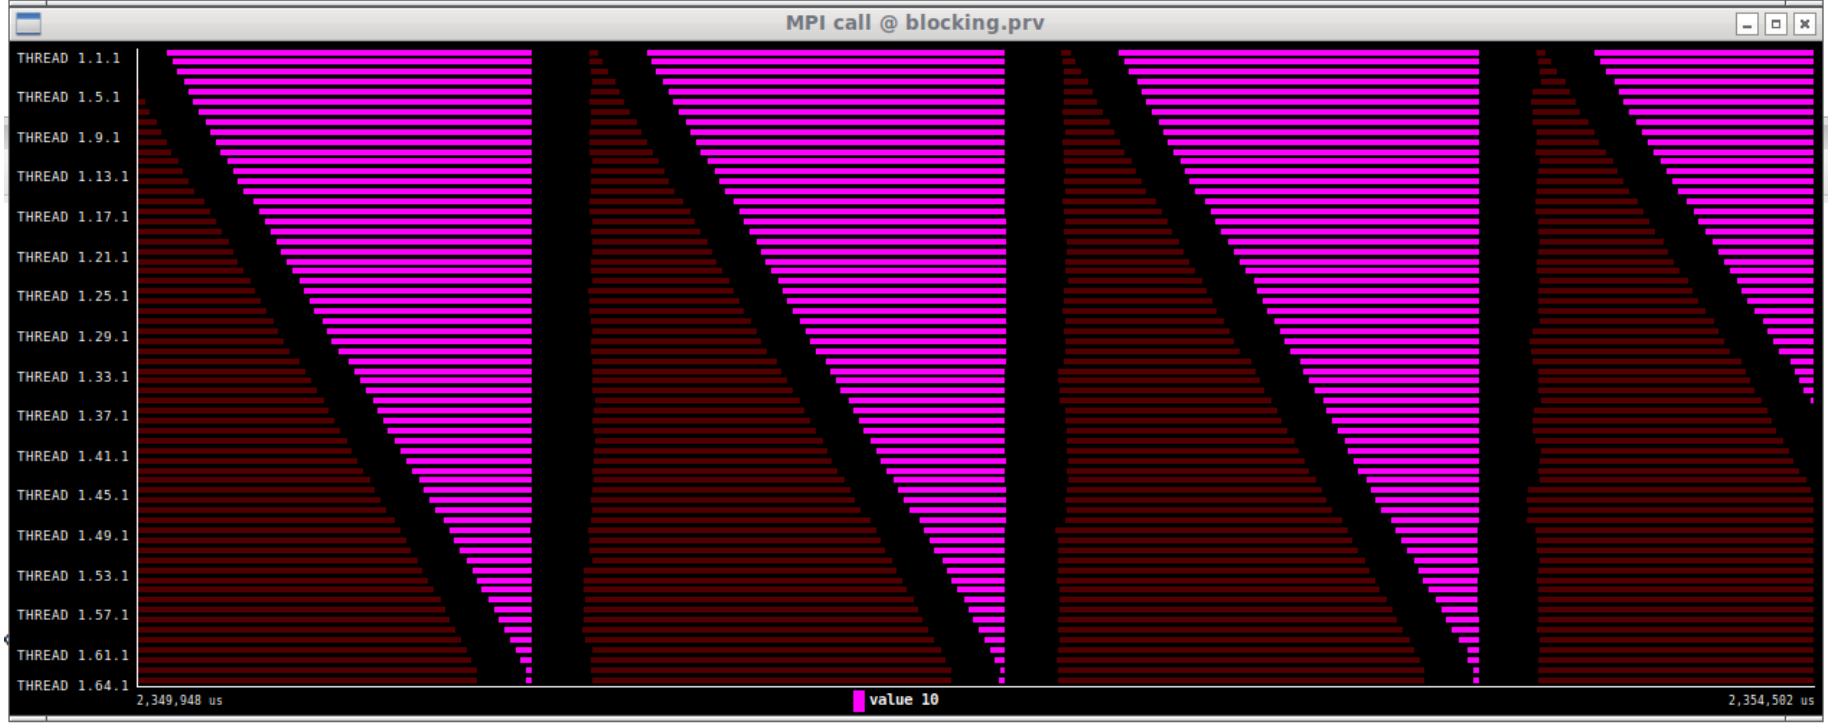
MPI call @ blocking.prv
THREAD 1.1.1
THREAD 1.5.1
THREAD 1.9.1
THREAD 1.13.1
THREAD 1.17.1
THREAD 1.21.1
THREAD 1.25.1
THREAD 1.29.1
THREAD 1.33.1
THREAD 1.37.1
THREAD 1.41.1
THREAD 1.45.1
THREAD 1.49.1
THREAD 1.53.1
THREAD 1.57.1
THREAD 1.61.1
THREAD 1.64.1
2,349,948 us
value 10
2,354,502 us

Non-blocking vs overlapping

	Outside MPI	MPI_Isend	MPI_Irecv	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank	MPI_Comm_size	MPI_Init	MPI_Finalize
THREAD 1.1.1	66.08 %	0.42 %	0.39 %	1.41 %	29.49 %	0.00 %	0.05 %	0.01 %	2.16 %
THREAD 1.2.1	41.85 %	0.72 %	0.57 %	1.35 %	29.23 %	0.00 %	0.04 %	25.44 %	0.79 %
THREAD 1.3.1	41.87 %	0.73 %	0.57 %	1.28 %	29.26 %	0.00 %	0.05 %	25.43 %	0.80 %
THREAD 1.4.1	41.71 %	0.73 %	0.58 %	1.39 %	29.32 %	0.00 %	0.03 %	25.42 %	0.81 %
THREAD 1.93.1	62.93 %	0.70 %	0.56 %	1.24 %	31.20 %	0.00 %	0.03 %	1.33 %	2.00 %
THREAD 1.94.1	63.19 %	0.71 %	0.57 %	1.23 %	31.27 %	0.00 %	0.03 %	0.96 %	2.03 %
THREAD 1.95.1	63.57 %	0.71 %	0.56 %	1.18 %	31.30 %	0.00 %	0.04 %	0.59 %	2.04 %
THREAD 1.96.1	63.74 %	0.38 %	0.32 %	1.79 %	31.45 %	0.00 %	0.04 %	0.22 %	2.06 %

Non-blocking parallel version

	Outside MPI	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank	MPI_Comm_size	MPI_Init	MPI_Finalize	MPI_Iallreduce
THREAD 1.1.1	66.71 %	0.29 %	0.41 %	28.34 %	1.55 %	0.02 %	0.00 %	0.04 %	0.01 %	2.17 %	0.46 %
THREAD 1.2.1	43.76 %	0.54 %	0.67 %	27.75 %	1.44 %	0.11 %	0.00 %	0.03 %	24.46 %	0.71 %	0.53 %
THREAD 1.3.1	43.50 %	0.58 %	0.55 %	28.10 %	1.52 %	0.12 %	0.00 %	0.03 %	24.45 %	0.72 %	0.43 %
THREAD 1.4.1	43.58 %	0.55 %	0.66 %	27.66 %	1.69 %	0.12 %	0.00 %	0.03 %	24.44 %	0.73 %	0.54 %
THREAD 1.93.1	63.58 %	0.59 %	0.60 %	29.85 %	1.35 %	0.14 %	0.00 %	0.03 %	1.32 %	2.01 %	0.53 %
THREAD 1.94.1	63.93 %	0.59 %	0.60 %	29.87 %	1.32 %	0.14 %	0.00 %	0.03 %	0.95 %	2.03 %	0.53 %
THREAD 1.95.1	64.20 %	0.59 %	0.60 %	29.97 %	1.30 %	0.14 %	0.00 %	0.04 %	0.59 %	2.05 %	0.52 %
THREAD 1.96.1	64.36 %	0.29 %	0.31 %	30.13 %	1.90 %	0.14 %	0.00 %	0.04 %	0.23 %	2.08 %	0.52 %

Overlapping parallel version

epcc

THE UNIVERSITY OF EDINBURGH

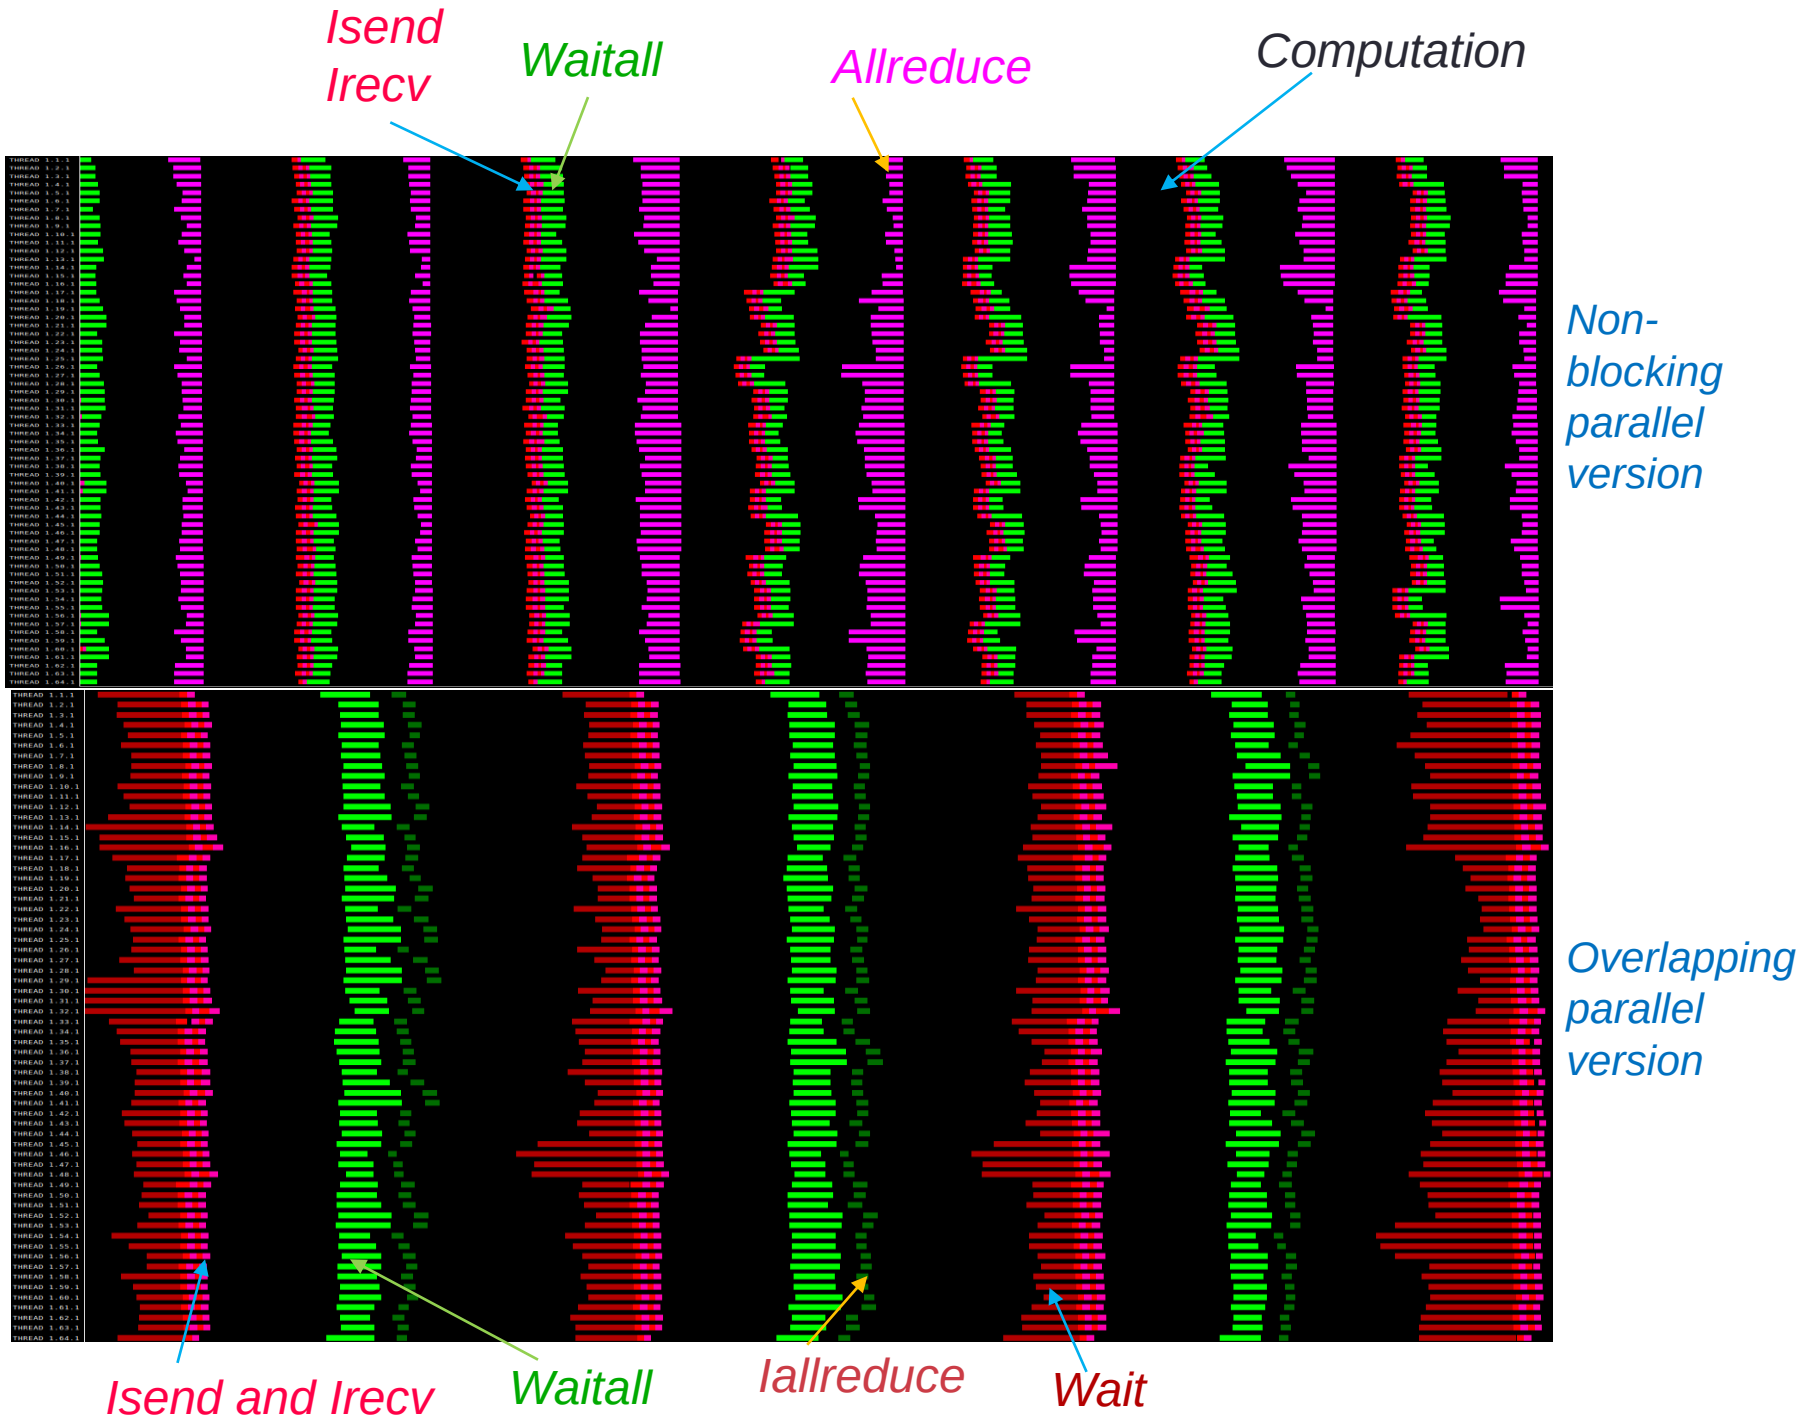
Isend
Irecv
Waitall
Allreduce
Computation
Non-blocking parallel version
Overlapping parallel version
Isend and Irecv
Waitall
Iallreduce
Wait