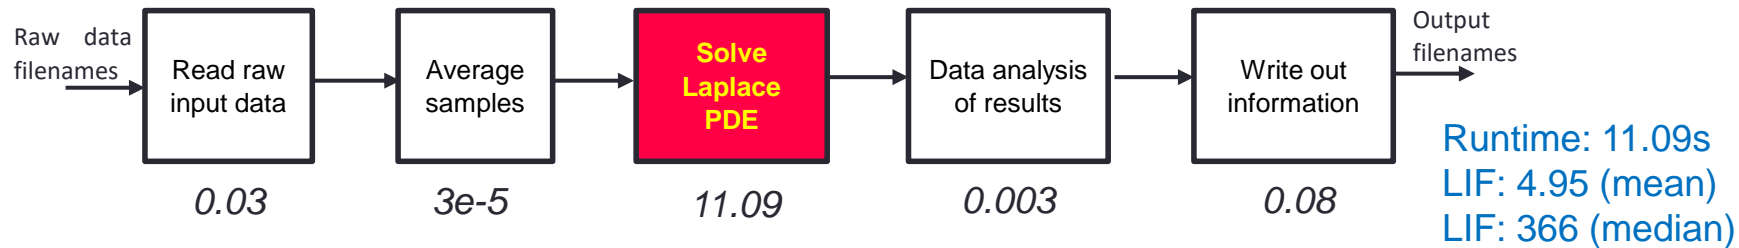


Practical three

Wrap up


Load imbalance factor

- Sample solutions are available
 - MPI P2P messages for communication between stages
 - For the termination poisoned pill an empty (NULL) message can be sent

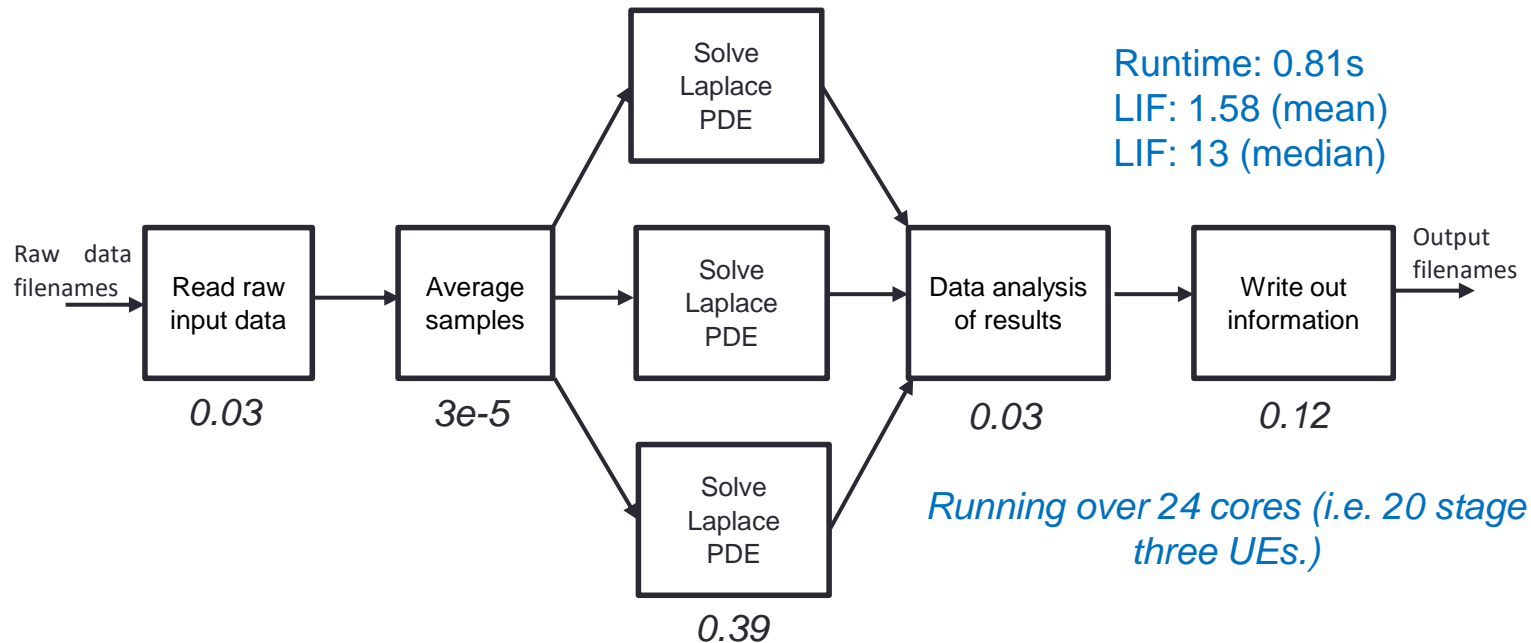


- But the stages of the pipeline are heavily imbalanced
 - Fact that we have one very large number and the rest small numbers impacts how we define the term *average*
 - Not necessarily easy to give lightly loaded stages more work, but can do something to optimise the heavily loaded stage(s)

Let's look at this with Paraver




Duplicating the third stage



- All extra UEs make up duplicate stage three.
 - No stage three UEs communicate, but instead work concurrently on different pieces of data
- Fairly simple to do, but termination does require a little more thought

In Paraver



A note on MPI P2P communication calls

- Blocking calls
 - MPI_Send, MPI_Ssend, MPI_Bsend, MPI_Rsend
 - Will only return from the call when the communication has *completed*
 - The data buffer can be reused
- Non-blocking calls
 - MPI_Isend, MPI_Issend, MPI_Ibsend, MPI_Irsend
 - Does not wait until the communication completes, instead returns pretty much immediately and need to check request handle to track completion
 - Via MPI_Wait, MPI_Test, MPI_Cancel etc....

Irrespective, can not reuse data buffer until P2P communication call has completed!

But what does “completed” mean?

- Depends on whether it's a send, send, bsend, rsend
- Standard send (MPI_Send)
 - Completes either when the message has been copied to an internal buffer or started to be received by the receiver
- Buffered send (MPI_Bsend)
 - Completes when the message has been copied to an internal buffer
- Synchronous send (MPI_Ssend)
 - Completes when the message has started to be received by the receiver
- Ready send (MPI_Rsend)
 - Same as a standard send, but the receiver must have already posted a corresponding receive call, otherwise is erroneous
 - Suggest not so useful nowadays, but could improve performance historically

My suggestions

- Realise that blocking vs non-blocking is entirely separate than the semantics of completion
 - For MPI send calls, it's simpler for receives!
- Lots of people find this a difficult part of the standard to understand, so when looking at a user's parallel code can be a source of bugs/performance issues
- If you have lots of synchronous sends (MPI_Ssend) in your code, then do you really need them?
 - If you replace with standard send will this maintain correctness, as by doing so could give you a performance boost